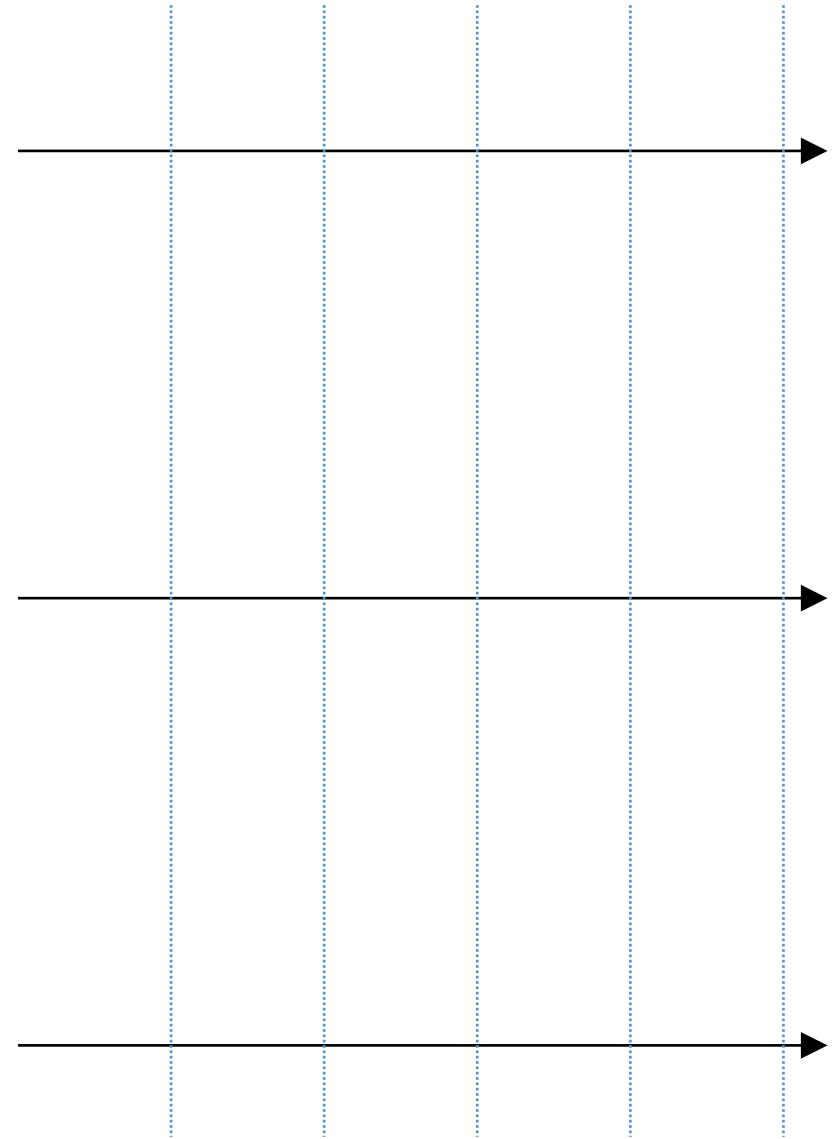
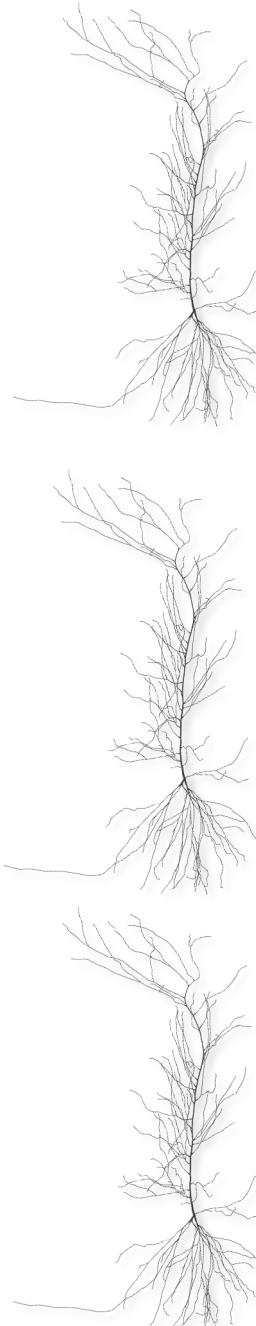


Building, running, and visualizing parallel NEURON models



Sync time points

Why write parallel code?



Get results faster.



Solve larger problems in the same amount of time.



Work on large problems that do not fit in a single machine's memory.



Run more simulations (e.g., parameter sweeps)



Move forward while waiting for data (e.g., from a networked machine.)

Why to **not** write parallel code

- Parallel code is more difficult to write and to debug.
- There are new types of bugs that only occur in parallel contexts. (e.g., race conditions).
- It is even possible to parallelize code and end up with something that runs slower than the original (e.g., due to communication overhead).
- Your time matters too; will the run-time savings be more than the extra time that you spend writing?

Three main classes of parallel problems

Parameter sweeps

- Running the same (typically fast) simulation 1000s of times with different parameters is an example of an *embarrassingly parallel problem*. This can be done without MPI using Python's built-in multiprocessing module, or with MPI using NEURON's bulletin boards.

Distributing networks across processors

- Cells can communicate via:
 - Logical spike events with significant axonal or synaptic delay.
 - Postsynaptic conductance depending continuously on presynaptic voltage.
 - Gap junctions.

Distributing single cells across processors

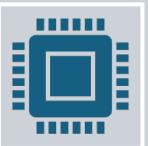
- The *multisplit* method distributes portions of the tree cable equation across different machines.

A parallel model can fall in any or all these classes.

Some parallel philosophy



A network of neurons is composed of many individual neurons of potentially many cell types. As much as possible, design and debug each cell type separately before building the network.



A simulation should give the same results regardless of the number of processors used to run it.



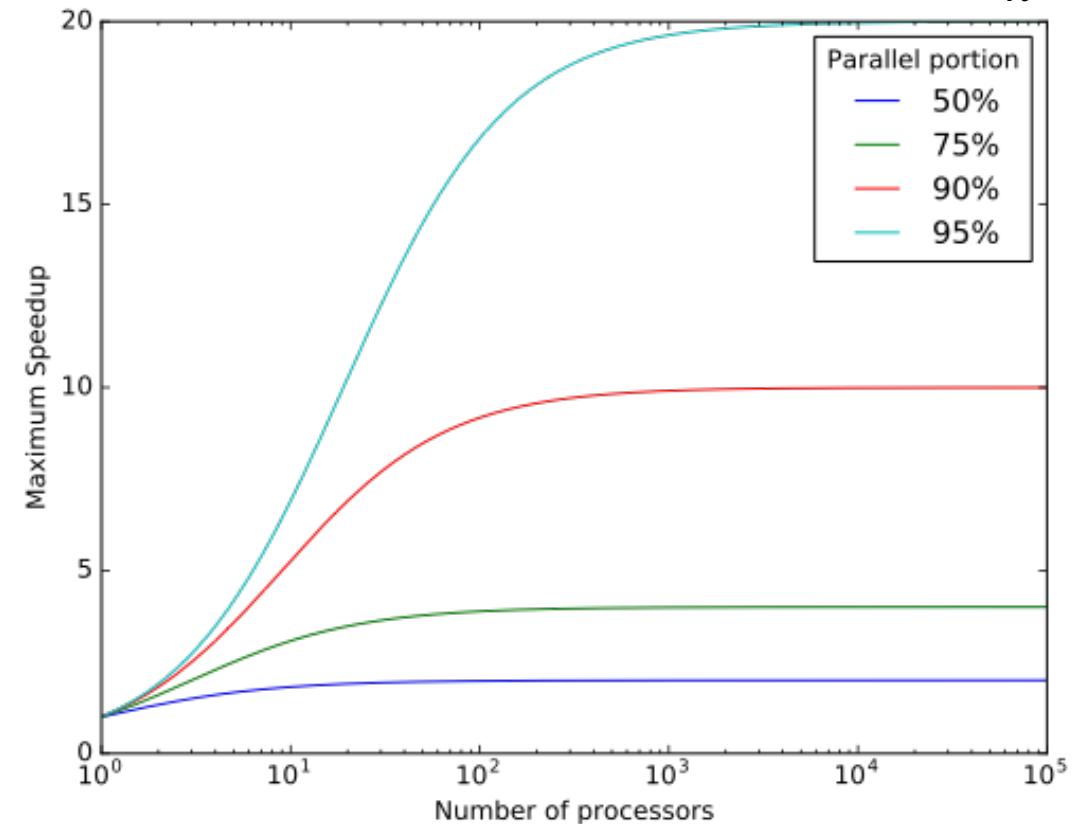
When possible, parameterize your network so you can run a small test first.

A tiny bit of theory...

Parallel scaling performance

- Suppose I have a serial simulation that runs in 100 minutes. How fast could I expect a parallel version of it to run on 1000 processors?
- Suppose further that 1% of the code could not be parallelized. How does that affect your answer?

$$\text{Amdahl's Law: Max speedup} = \frac{1}{(1-p) + \frac{p}{n}}$$

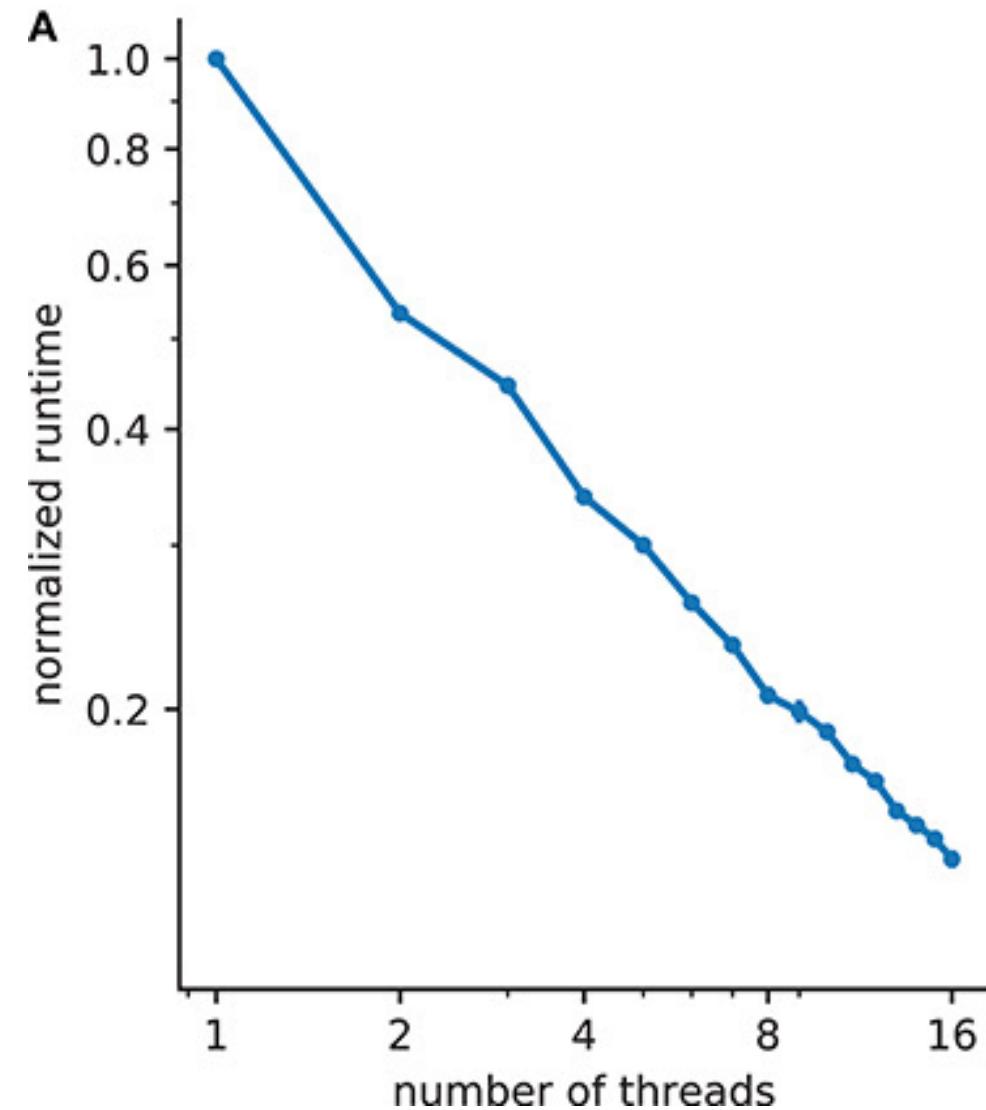


Implicit assumption

- Amdahl's law assumes that p is a constant.
- In practice, we probably do not want to do the same simulation in parallel that we did in serial; we probably want to do a bigger model.
- Loading e.g. neuron takes the same amount of time regardless of the size of our problem. So p may increase as n increases.

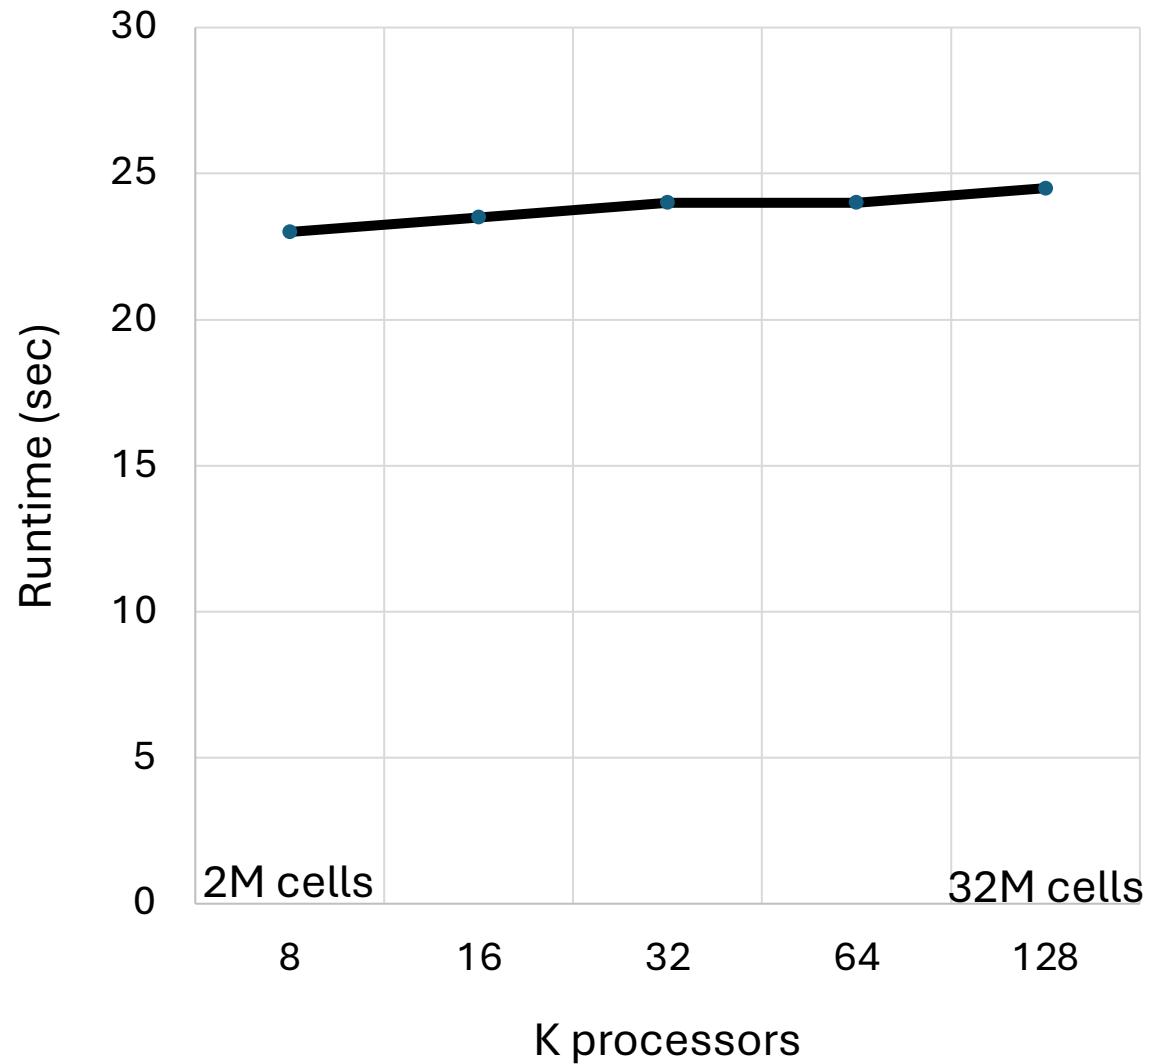
Strong scaling

- Fixed problem
- Measure speedup per processor
- On a log-log graph of threads/processes vs runtime, ideal is a graph that looks linear with slope -1



Weak scaling

- Fixed amount of work per thread/processor.
- Measure relationship between size of problem and runtime.
- Ideal is a flat relationship.



So how do we design a model to be parallelizable?

Who does what?

pc.nhost (size) and pc.id (rank)

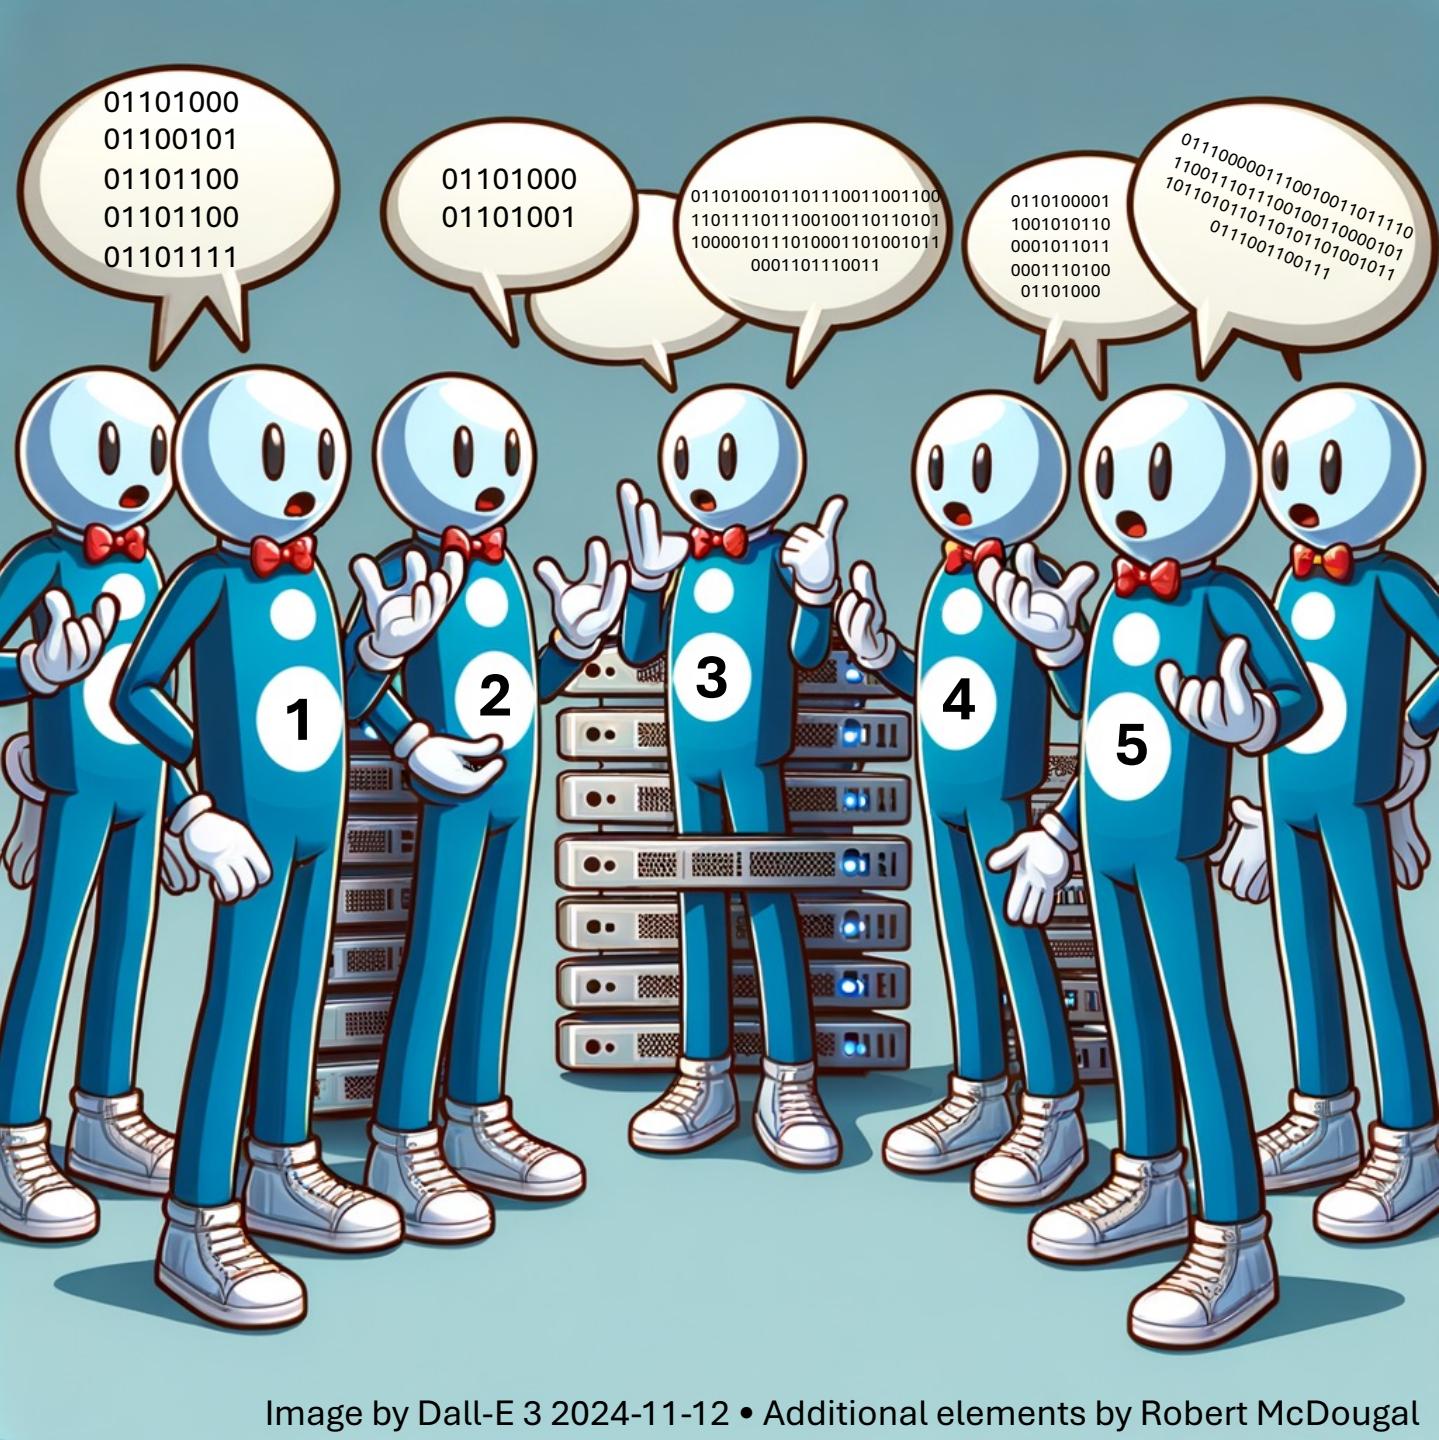
Idea: Use a predefined rule for splitting up the simulation. Since all processors know the rule, they know which parts of the simulation they must setup and simulate.

```
from neuron import h
h.nrnmpi_init()
pc = h.ParallelContext()

val = pc.id()
squared = val ** 2

print(
    f"I am rank {val} (of {pc.nhost()}) "
    f"and my square is {squared}"
)
```

```
% mpiexec -n 4 python squaretest.py
numprocs=4
I am rank 1 (of 4) and my square is 1
I am rank 2 (of 4) and my square is 4
I am rank 3 (of 4) and my square is 9
I am rank 0 (of 4) and my square is 0
```



That is, let's give all our cell classes a `gid`. If we parallelize, that will help us decide which processor id handles it.

- In addition, it will be convenient to specify morphology in a dedicated method and add a `__repr__` method to identify the object.
- Here, the `gid` should be a globally unique identifying integer. We do not use class variables to generate the integer automatically because: (1) the numbers should not repeat between different processors, and (2) we may wish to recreate a single specific cell instead of the entire network.

```
from neuron import h
h.load_file('import3d.hoc')

class Pyramidal:
    def __init__(self, gid):
        self._gid = gid
        self._setup_morphology()
    def _setup_morphology(self):
        cell = h.Import3d_SWC_read()
        cell.input('c91662.swc')
        i3d = h.Import3d_GUI(cell, 0)
        i3d.instantiate(self)
    def __repr__(self):
        return f'Pyramidal[{self._gid}]'
```

Discretize, declare channels, set parameters

```
class Pyramidal:
    def __init__(self, gid):
        self._gid = gid
        self._setup_morphology()
        self._discretize()
        self._add_channels()
    def _setup_morphology(self):
        cell = h.Import3d_SWC_read()
        cell.input('c91662.swc')
        i3d = h.Import3d_GUI(cell, 0)
        i3d.instantiate(self)
    def __repr__(self):
        return f'Pyramidal[{self._gid}]'
    def _discretize(self, d_lambda=0.1, freq=100):
        h.load_file('stdlib.hoc')
        for sec in self.all:
            sec.nseg = 1 + 2 * math.ceil(
                (sec.L / (d_lambda * h.lambda_f(freq, sec=sec))) / 2)
    def _add_channels(self):
        h.hh.insert(self.soma)
        h.pas.insert(self.all)
        h.pas.uninsert(self.soma)
        for sec in self.axon + self.dend + self.apic:
            for seg in sec:
                seg.pas.g = 0.001
```



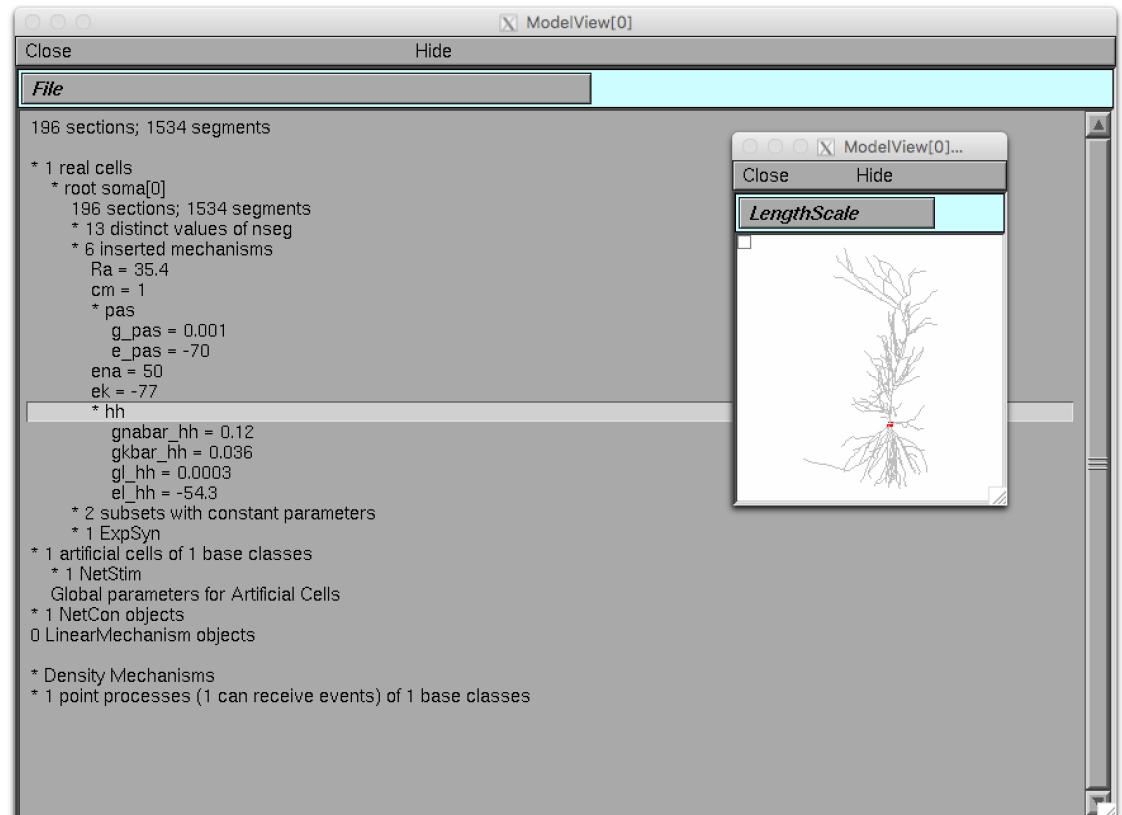
Python implementation of the d_lambda rule (Hines and Carnevale 2001).

Also: Ensures that there is an odd number of segments, which means there is a middle node. To test convergence, always multiply nseg by an odd number, e.g.,

```
for sec in self.all:
    sec.nseg *= 3
```

Examine for errors: Tools → ModelView

- Networks are made of cells.
- Always check each neuron or neuron type before combining into a network.
- Errors can still occur when neurons are combined into networks, but if the neurons are wrong then the network simulation definitely will be wrong.



Now: Run with pc.psolve(tstop)

```
from neuron import h
from matplotlib import pyplot
import math
h.load_file('stdrun.hoc')
h.load_file("import3d.hoc")
h.load_file("stdlib.hoc")

# class Pyramidal defined as before (omitted)

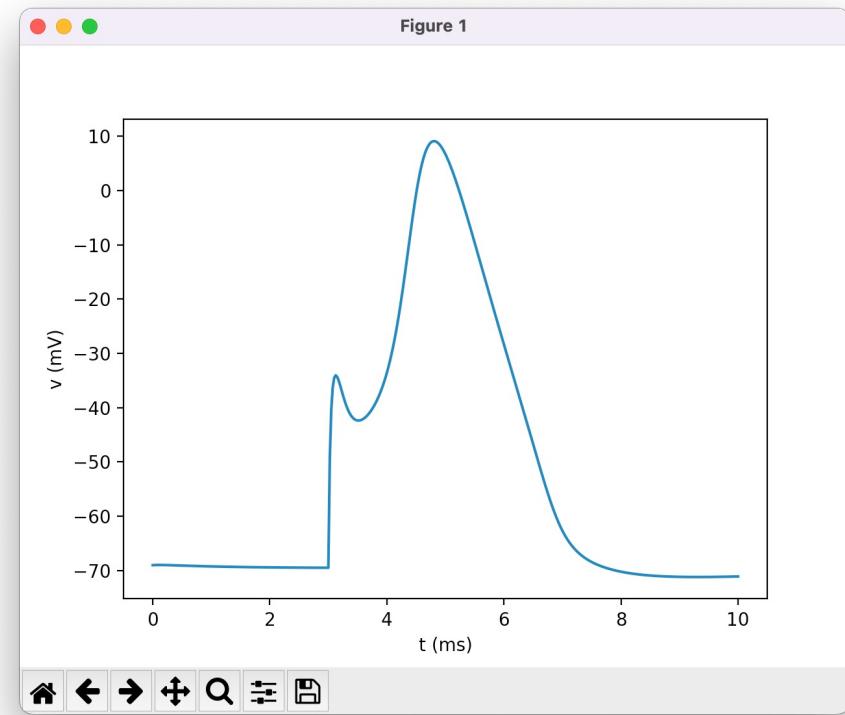
myPyramidal = Pyramidal(0)

postsyn = h.ExpSyn(myPyramidal.dend[0](0.5))
postsyn.e = 0 # reversal potential

stim = h.NetStim()
stim.number = 1
stim.start = 3
ncstim = h.NetCon(stim, postsyn)
ncstim.delay = 0
ncstim.weight[0] = 0.5

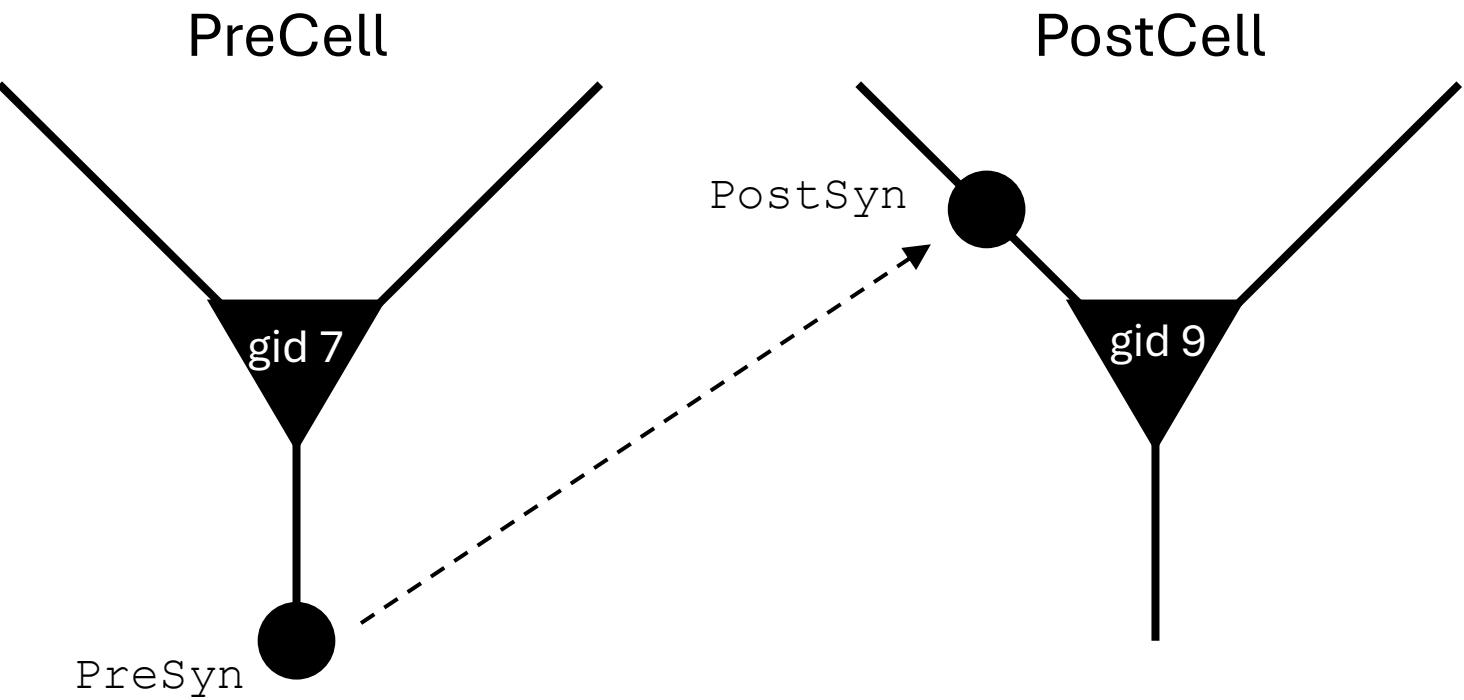
t = h.Vector().record(h._ref_t)
v = h.Vector().record(myPyramidal.soma[0](0.5)._ref_v)
pc = h.ParallelContext()
pc.set_maxstep(10)
h.finitialize(-69)
pc.psolve(10)
```

```
pyplot.plot(t, v)
pyplot.xlabel('t (ms)')
pyplot.ylabel('v (mV)')
pyplot.show()
```



Building synapses

- Typically, we may think of PreSyn as a voltage at a location that we want to monitor for crossing a threshold.
- PostSyn is some synaptic mechanism point process, e.g., `h.ExpSyn`.
- In a serial model, we would connect directly with an `h.NetCon`.



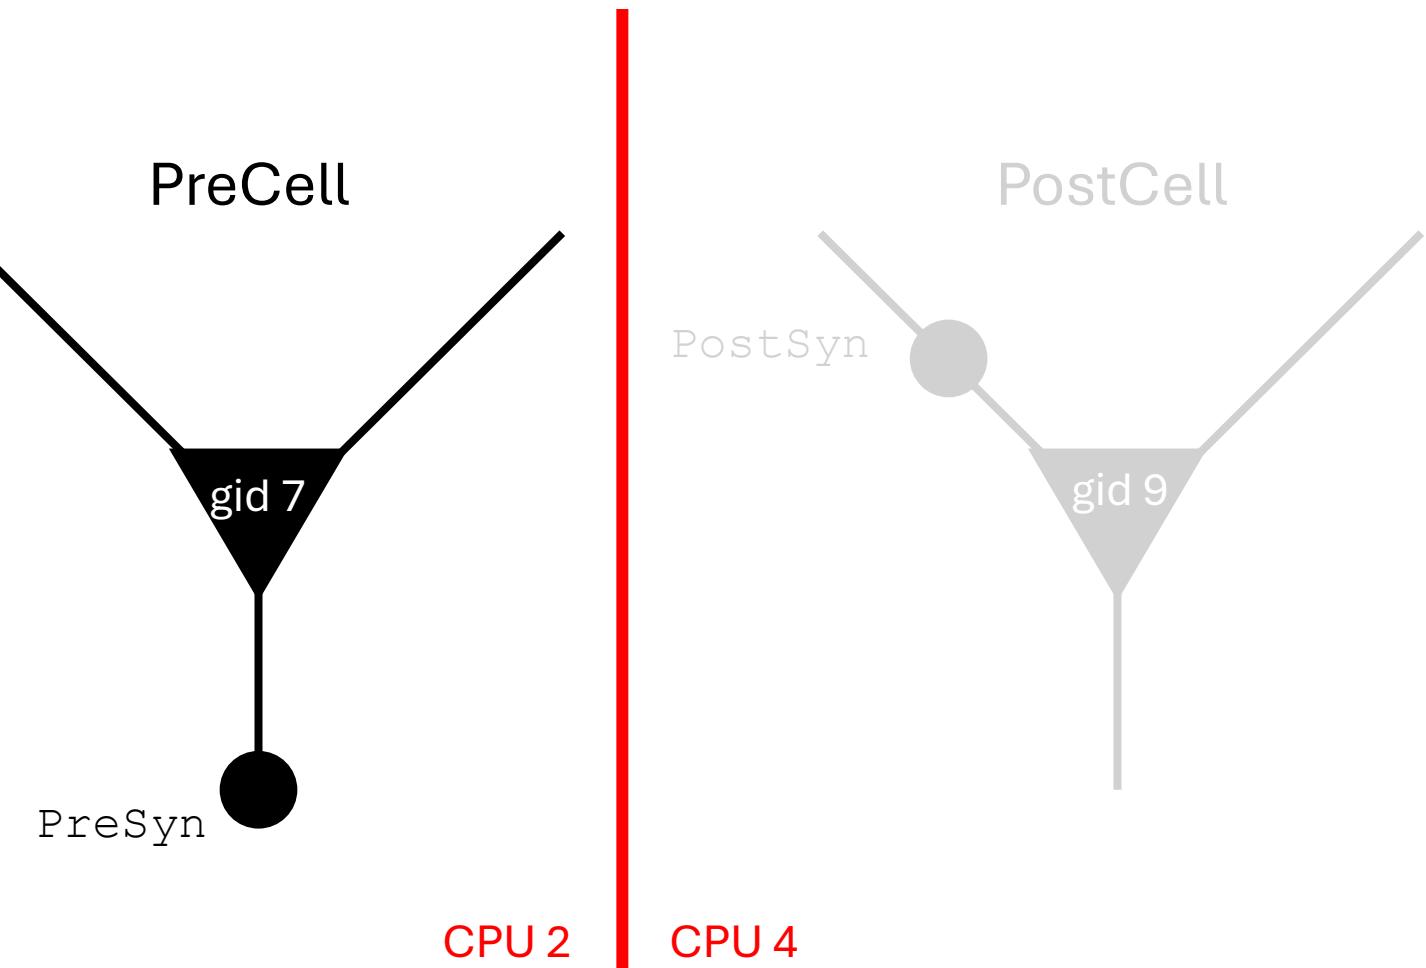
Configuring the presynaptic connection site

- In parallel contexts, one CPU cannot access variables stored on another CPU, so instead we must use the gids.
- Create cell only where the gid exists:

```
if pc.gid_exists(7):  
    PreCell = Cell()
```

- Associate gid with spike source:

```
nc = h.NetCon(PreSyn, None, sec=presec)  
pc.cell(7, nc)
```

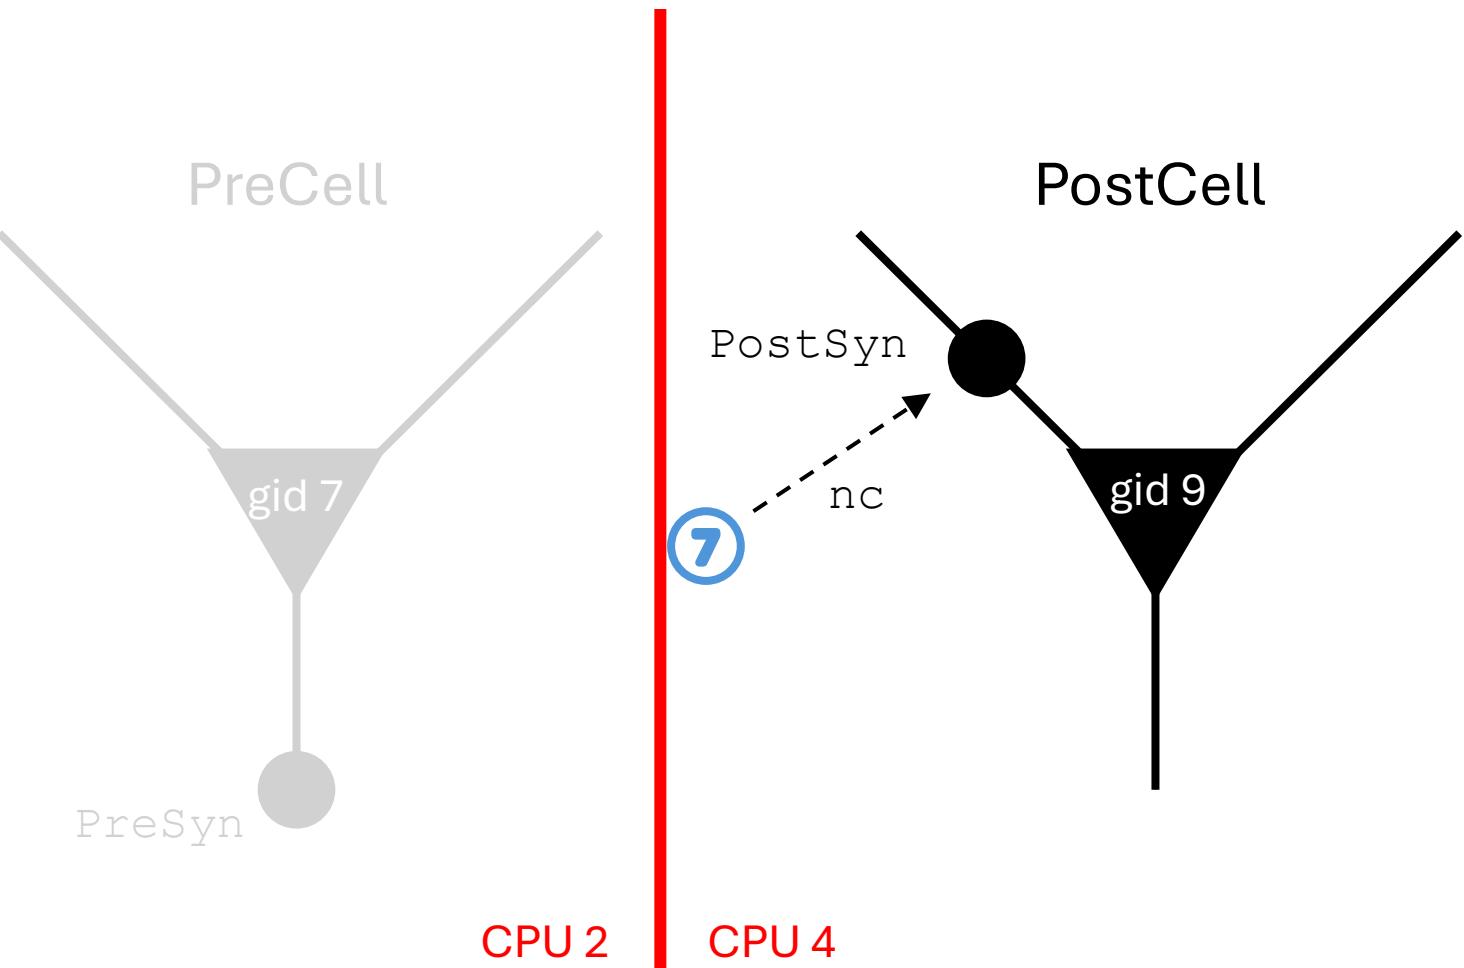


Configuring the postsynaptic connection site

- Create `h.NetCon` on node where target exists:

```
nc = pc.gid_connect(7, PostSyn)
```

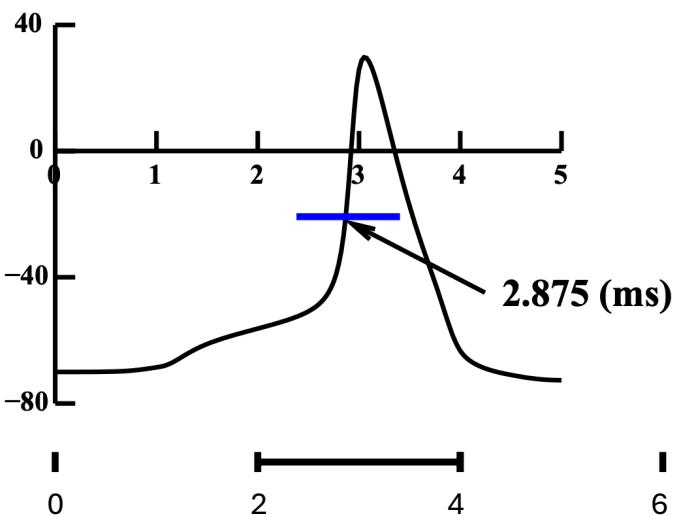
- PostSyn here is a Point Process with a `NET_RECEIVE` block, e.g., an `h.ExpSyn` but there are many synapse models available at <https://modeldb.science>



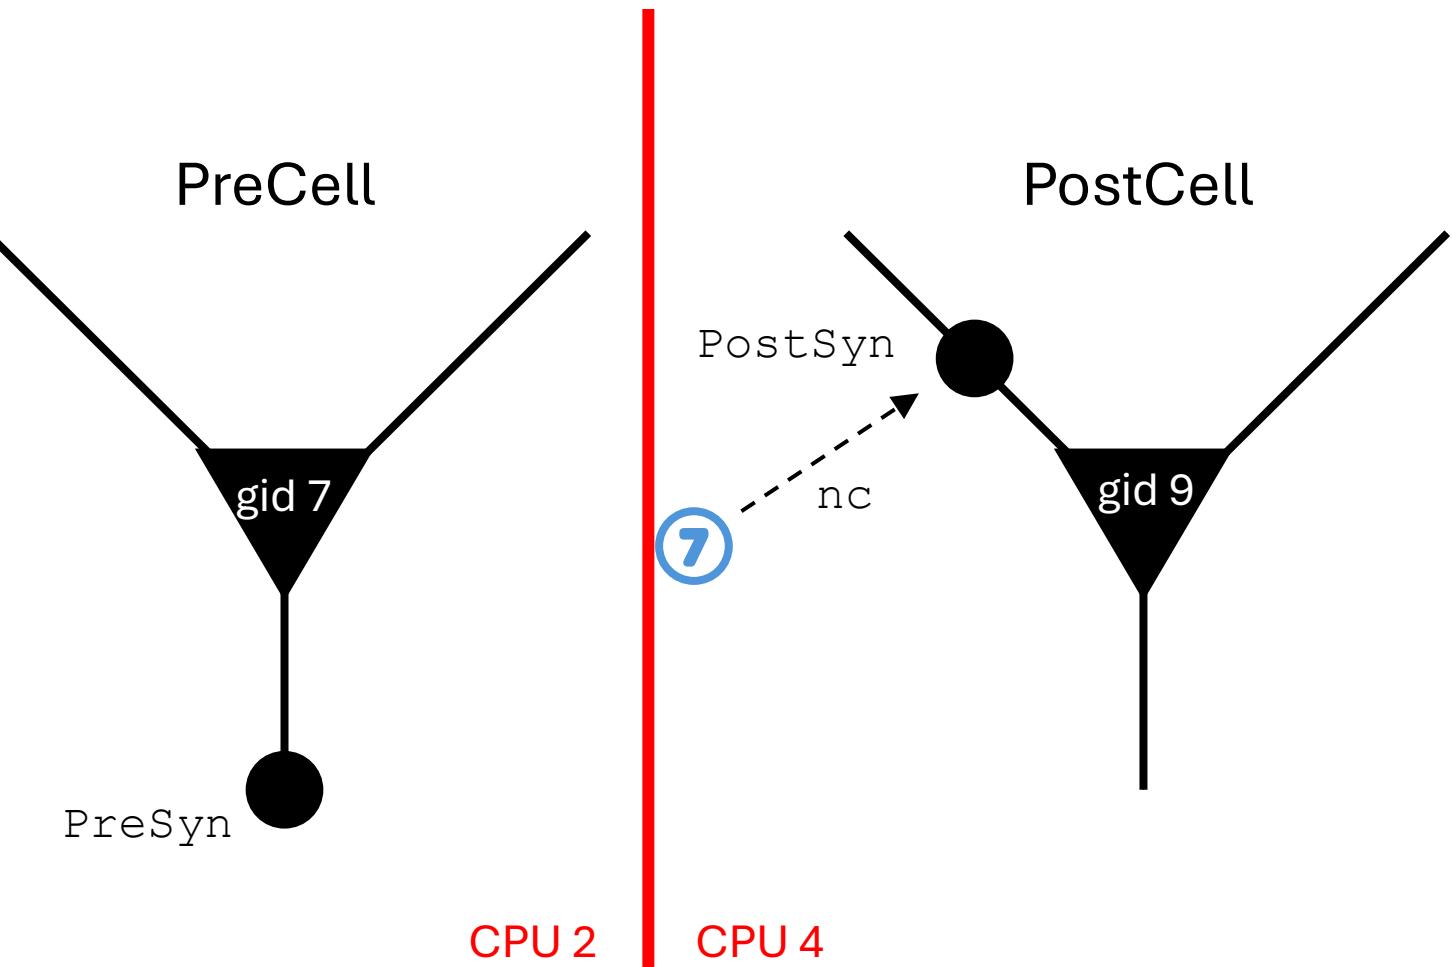
For this example, assume all synaptic delays are 2 ms.

Spike exchange method

- Step 1: Detect the time of a spike.



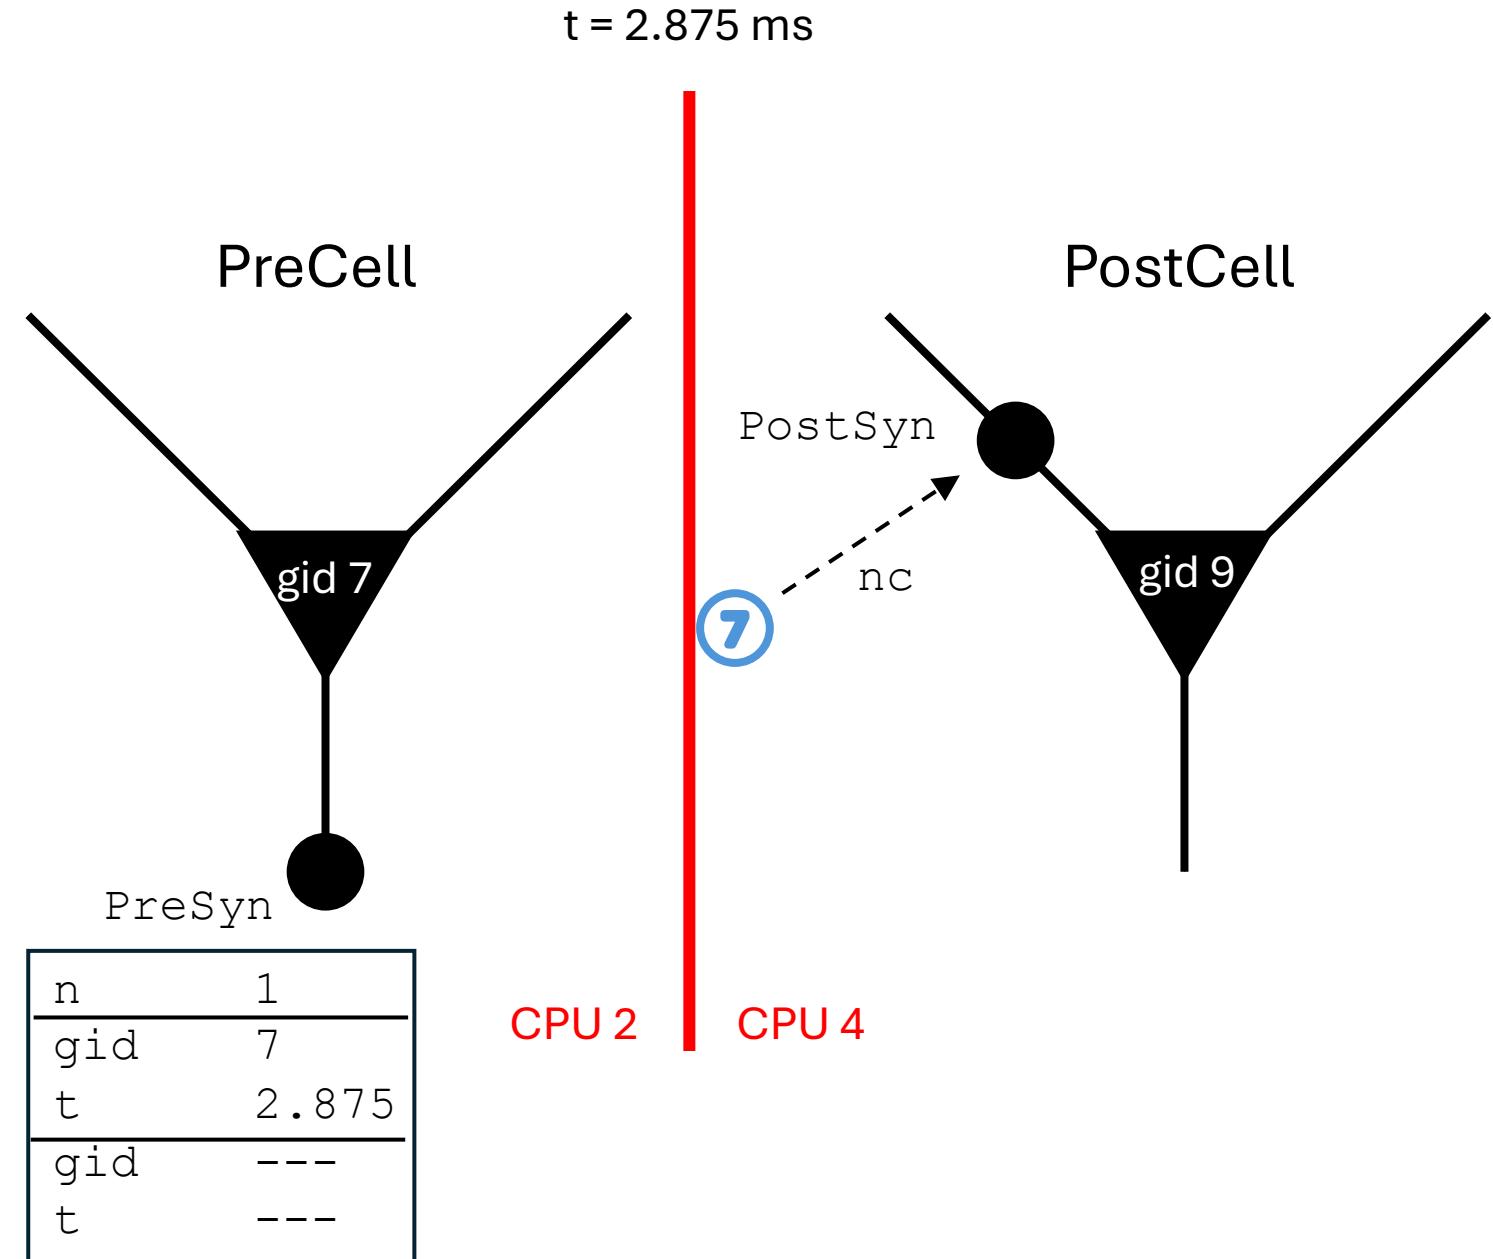
$t = 2.875 \text{ ms}$



For this example, assume all synaptic delays are 2 ms.

Spike exchange method

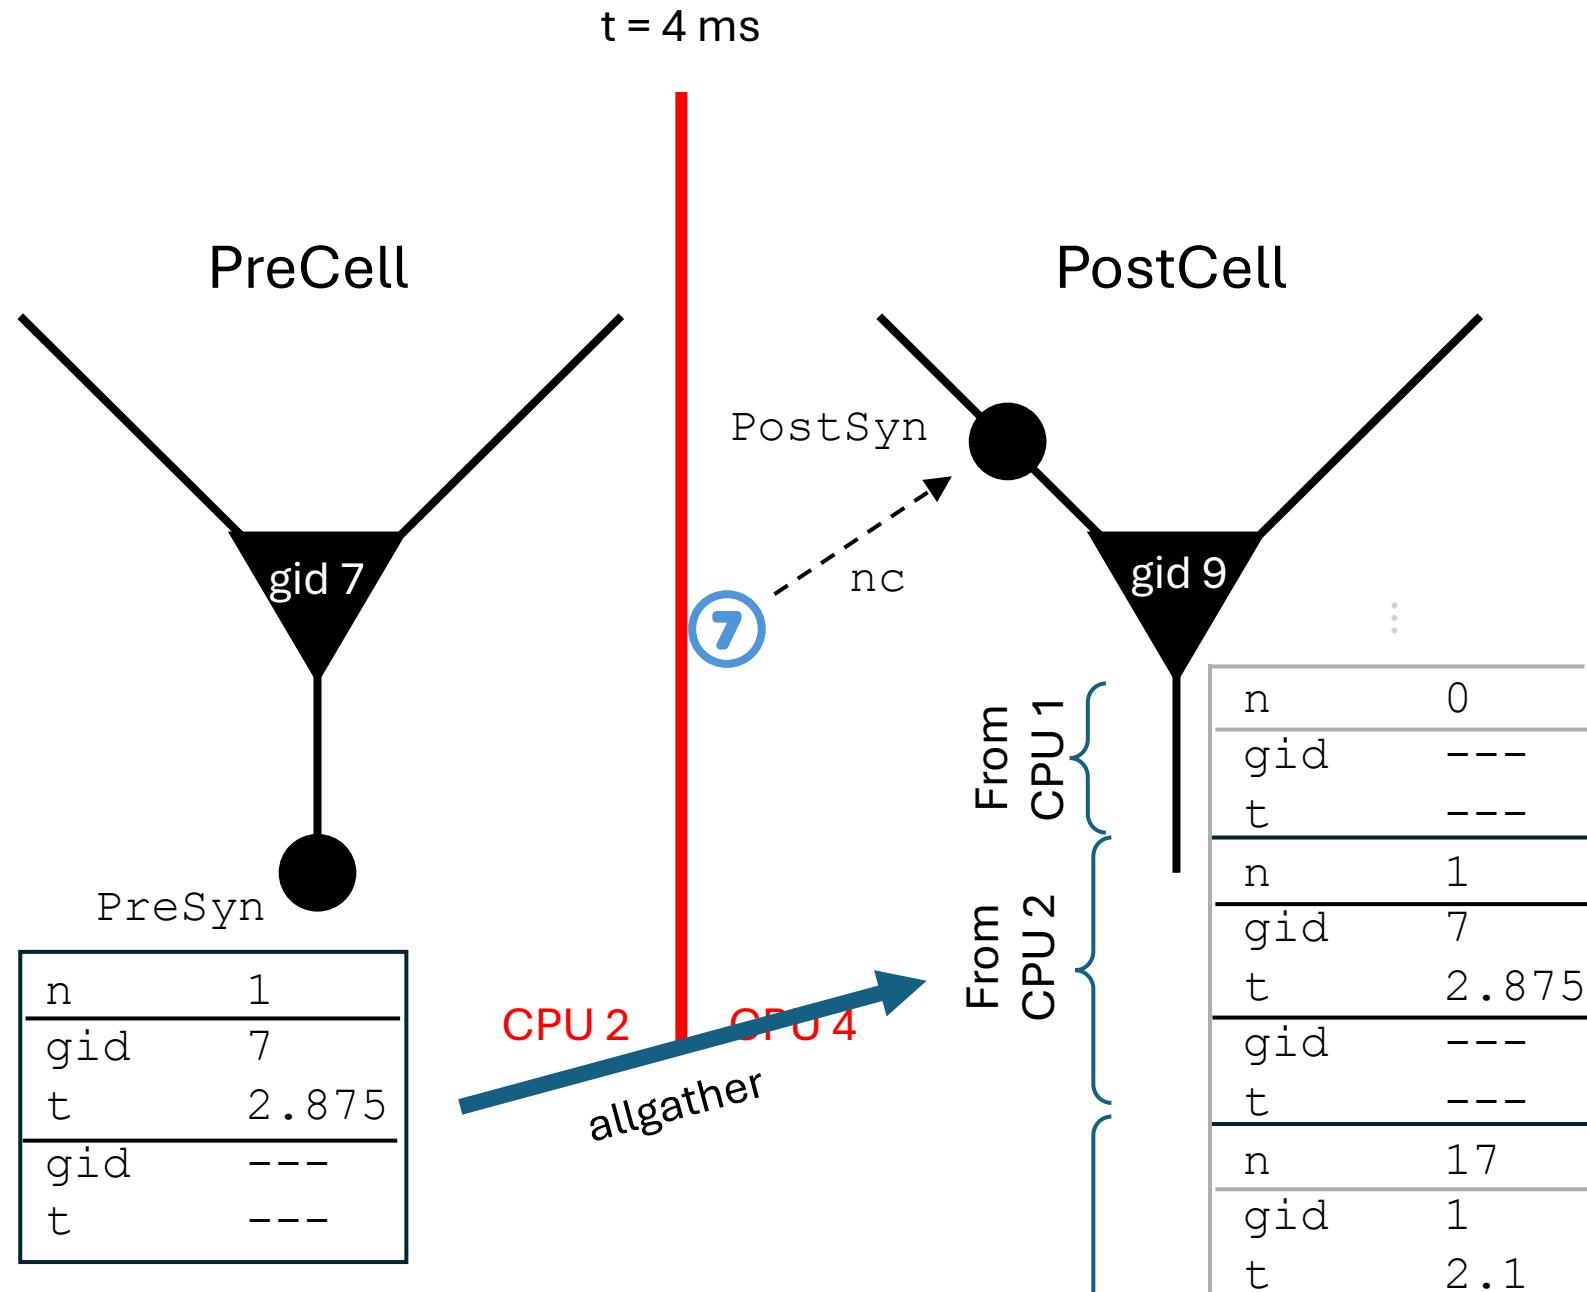
- Step 1: Detect the time of a spike.
- Step 2: Store in a queue.



For this example, assume all synaptic delays are 2 ms.

Spike exchange method

- Step 1: Detect the time of a spike.
 - Step 2: Store in a queue.
 - Step 3: Transfer via MPI's allgather.



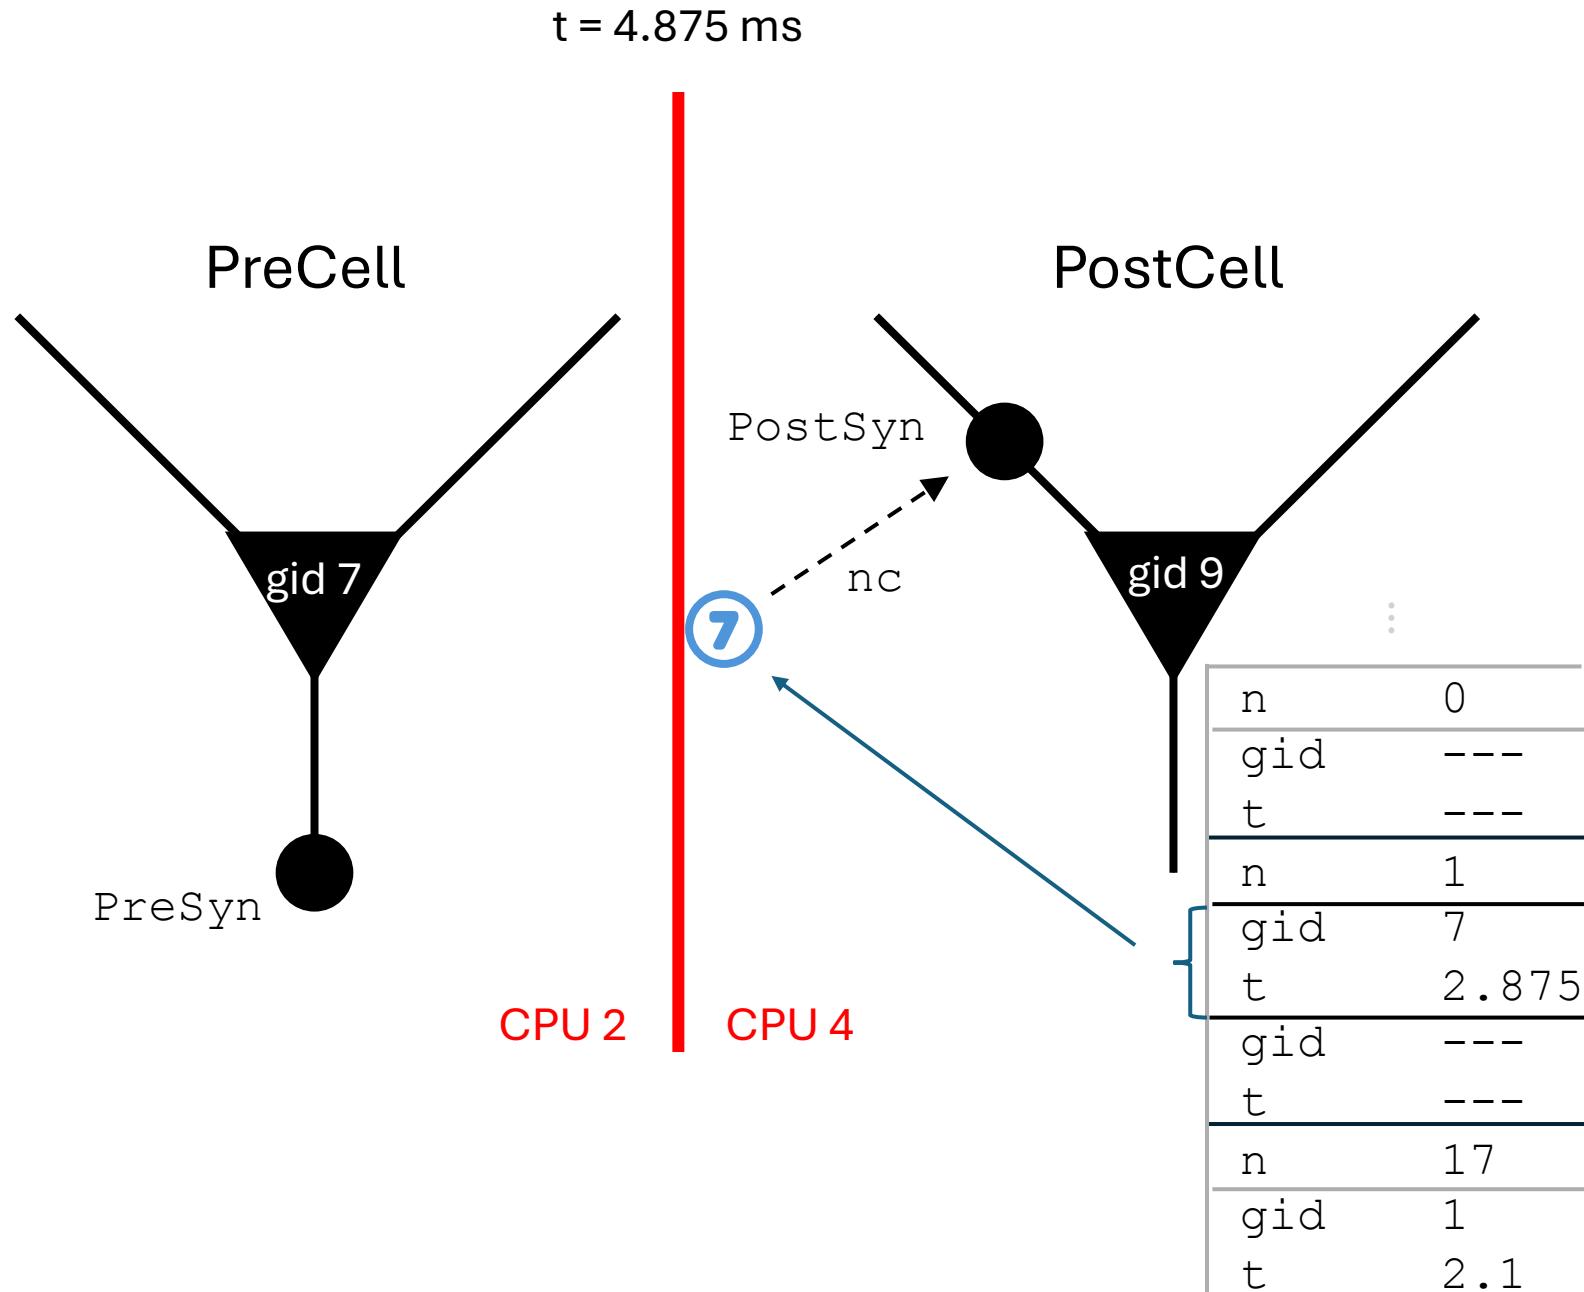
Collective communication: allgather



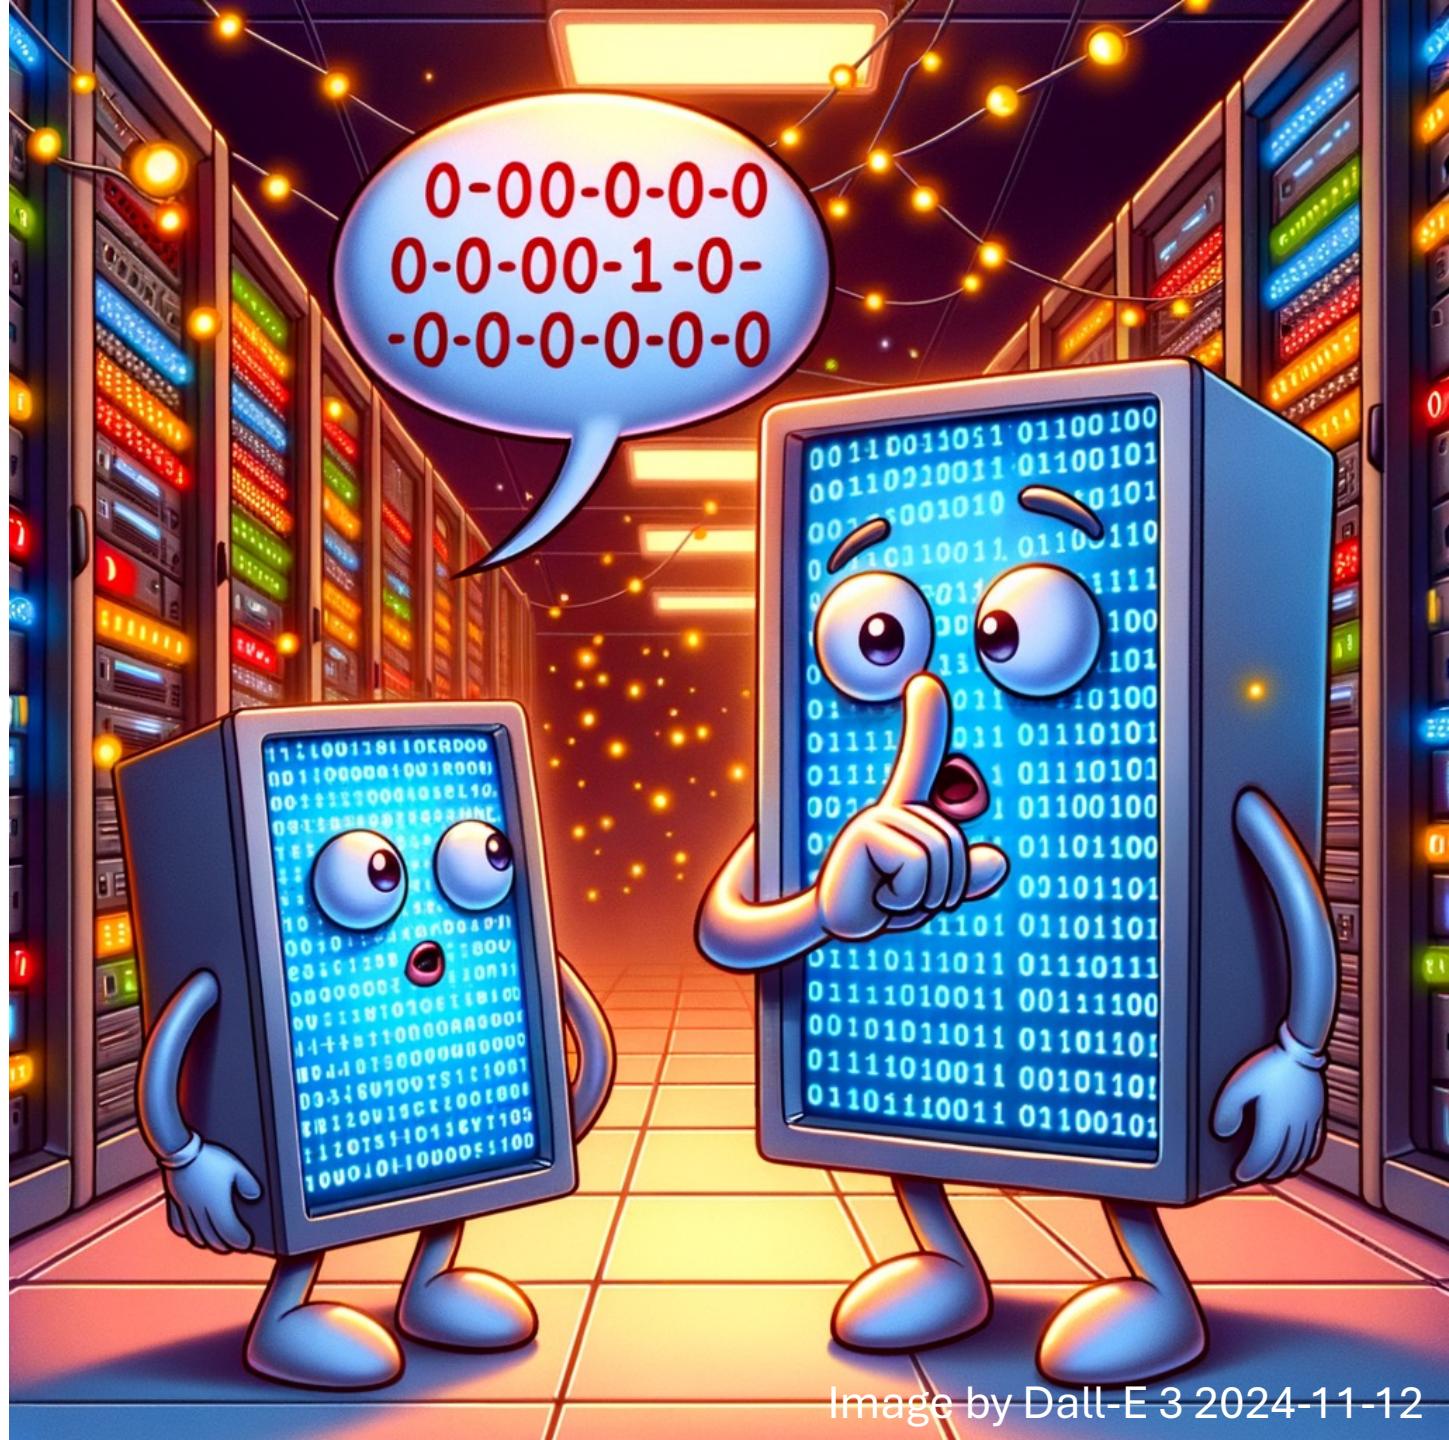
For this example, assume all synaptic delays are 2 ms.

Spike exchange method

- Step 1: Detect the time of a spike.
- Step 2: Store in a queue.
- Step 3: Transfer via MPI's allgather.
- Step 4: Deliver events at the appropriate time.



Minimize communication

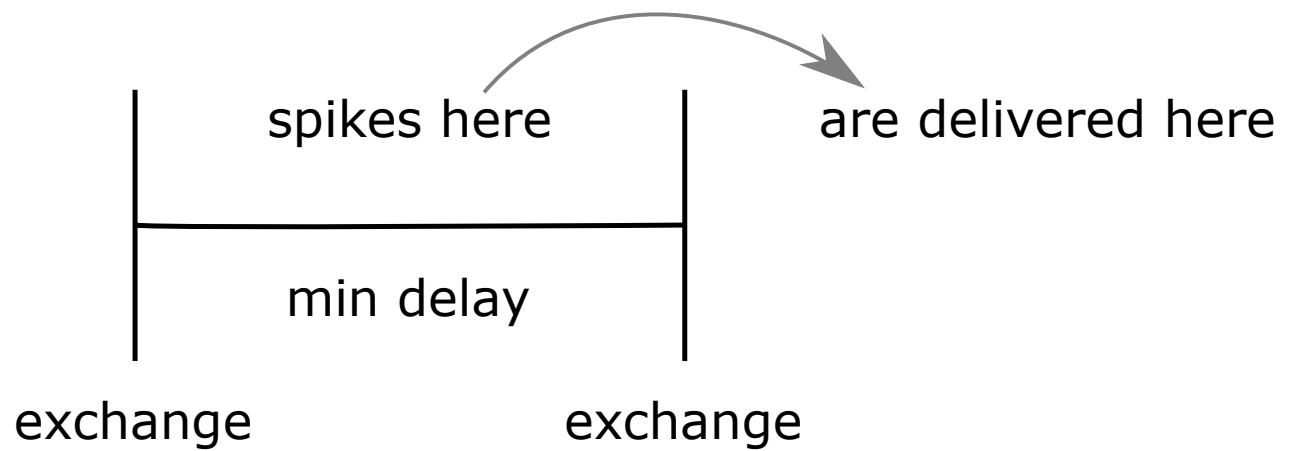


pc.set_maxstep

- NEURON will pick an event exchange interval equal to the smaller of all the `h.NetCon` delays and of the argument to `pc.set_maxstep`.
- In general, larger intervals are better because they reduce communication overhead.

Run parallel simulations using the idiom:

```
pc.set_maxstep(10)  
h.finitialize(-65)  
pc.psolve(tstop)
```



`pc.set_maxstep` must be called on each node; it uses MPI's `allreduce` to determine the minimum delay.

Back to our example...

- For most models, the delay due to axon propagation can be incorporated into a synaptic delay and thus it suffices to only make one connection point at the soma or axon hillock.
- `pc.set_gid2node` must be called before `pc.cell`.

```
class Pyramidal:  
    def __init__(self, gid):  
        self._gid = gid  
        self._setup_morphology()  
        self._discretize()  
        self._add_channels()  
        self._register_netcon()  
    def _register_netcon(self):  
        self.nc = h.NetCon(  
            self.soma[0](0.5)._ref_v, None,  
            sec=self.soma[0])  
        self.nc.threshold = 0  
        pc = h.ParallelContext()  
        pc.set_gid2node(self._gid, pc.id())  
        pc.cell(self._gid, self.nc)  
        # the rest of the class stays unchanged
```

Building a two cell network

Note: we use for loops and list comprehensions even when there is only two cells to avoid repeating ourselves (the DRY-principle) and to allow future generalization.

```
from neuron.units import ms, mV

class Network:
    def __init__(self):
        self.cells = [Pyramidal(i) for i in range(2)]
        # setup an exciteable ExpSyn on each cell's dendrites
        self.syns = [h.ExpSyn(cell.dend[0](0.5)) for cell in self.cells]
        for syn in self.syns:
            syn.e = 0 * mV
        # connect cell 0 to cell 1
        pc = h.ParallelContext()
        pre = 0
        post = 1
        self.nc = pc.gid_connect(pre, self.syns[post])
        self.nc.delay = 1 * ms
        self.nc.weight[0] = 0.5

n = Network()
```

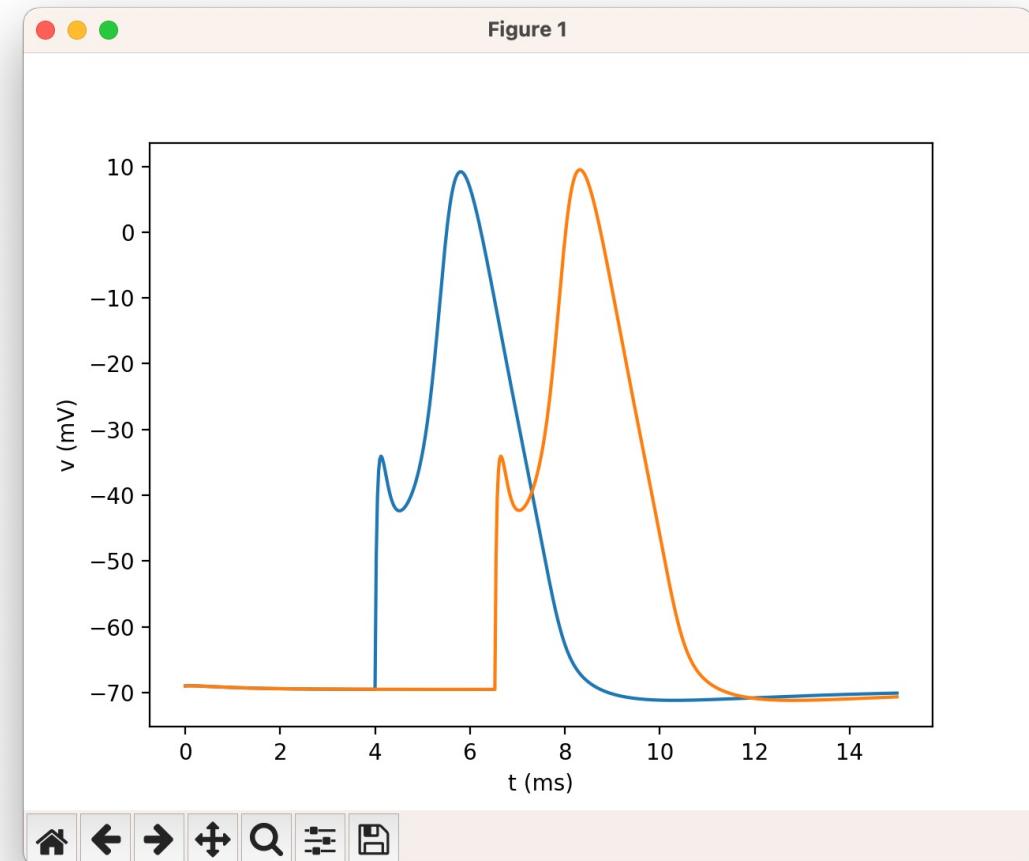
Running the two cell network

```
# drive the 0th cell
stim = h.NetStim()
stim.number = 1
stim.start = 3 * ms
ncstim = h.NetCon(stim, n.syns[0])
ncstim.delay = 1 * ms
ncstim.weight[0] = 0.5

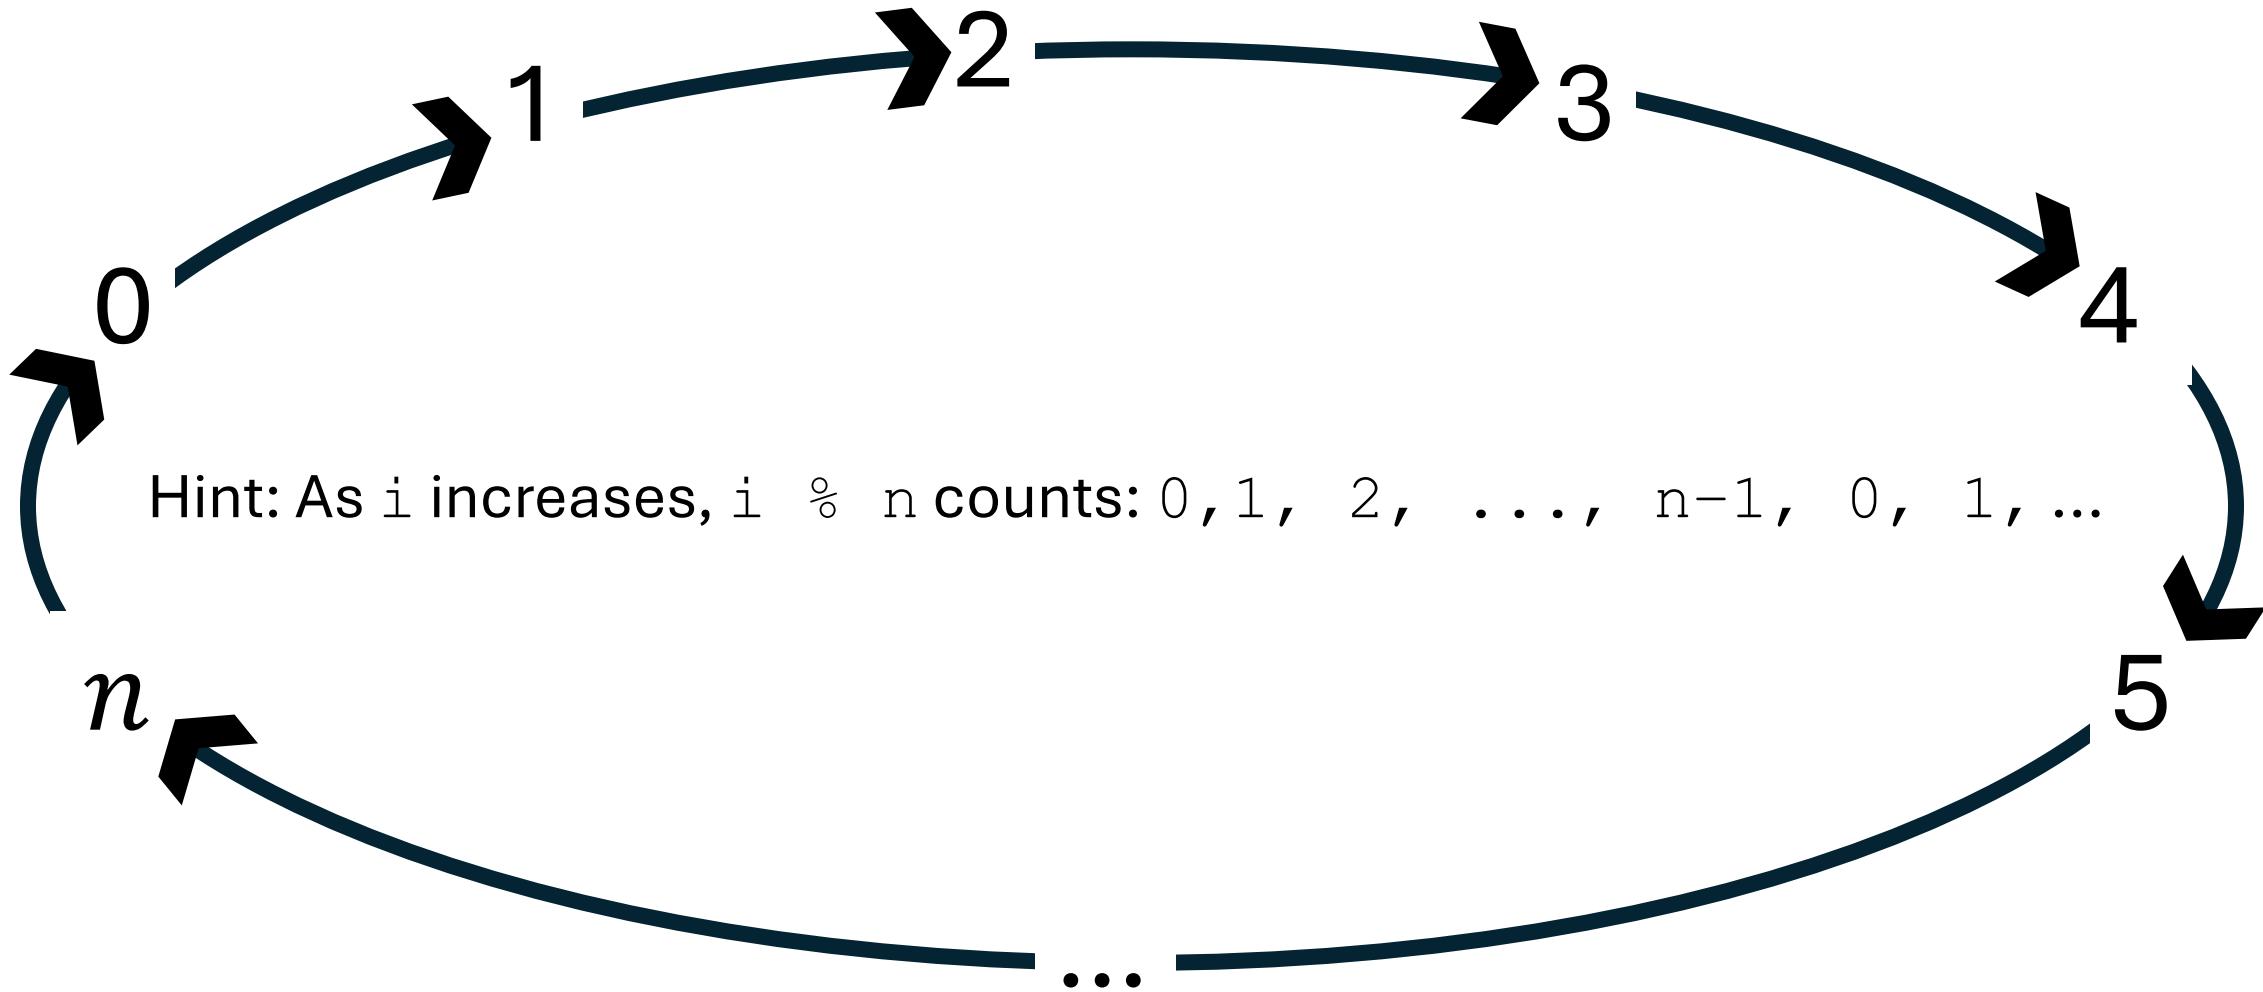
t = h.Vector().record(h._ref_t)
v = [h.Vector().record(cell.soma[0](0.5)._ref_v)
     for cell in n.cells]

pc = h.ParallelContext()
pc.set_maxstep(10 * ms)
h.finitiaize(-69 * mV)
pc.psolve(10 * ms)

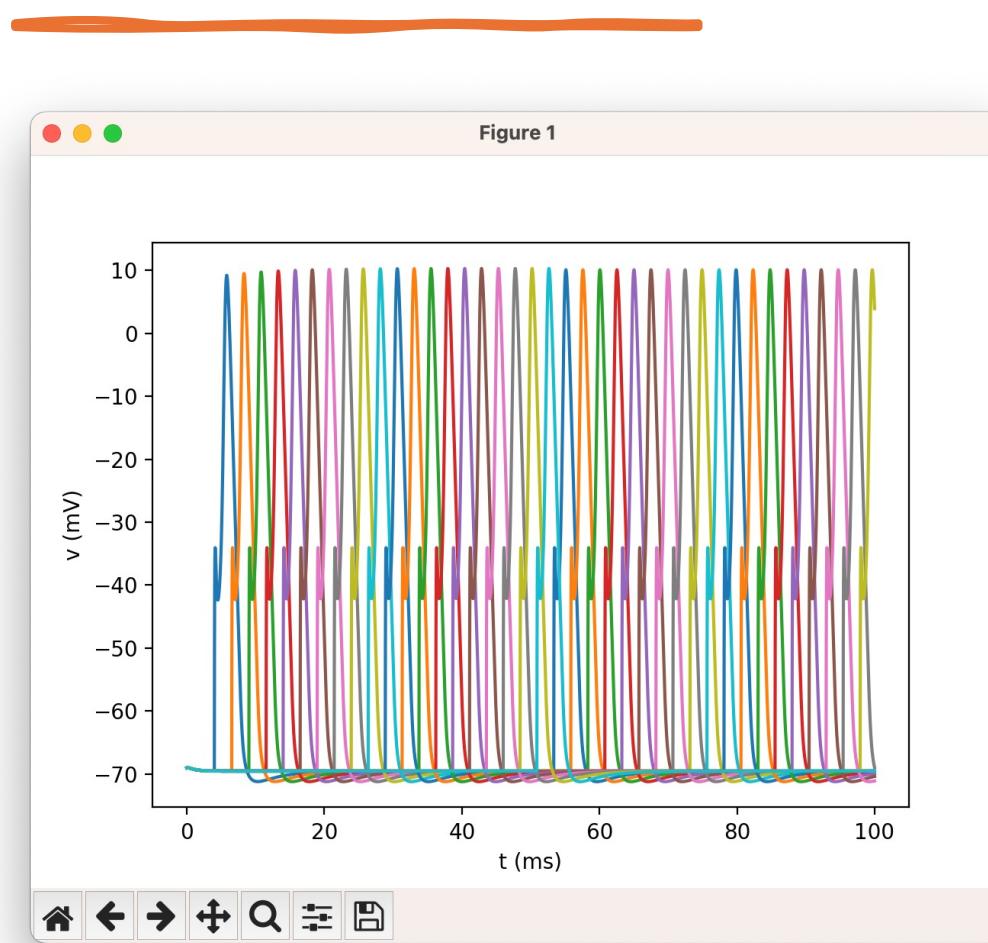
for myv in v:
    pyplot.plot(t / ms, myv / mV)
pyplot.xlabel('t (ms)')
pyplot.ylabel('v (mV)')
pyplot.show()
```



How do we generalize to a ring of n neurons?



Generalizing to n cells in a ring network (100 ms)



```
class Network:
    def __init__(self, num):
        self.cells = [Pyramidal(i) for i in range(num)]
        # setup an excitable ExpSyn on each cell's
        # dendrites
        self.syns = [h.ExpSyn(cell.dend[0](0.5)) for
                    cell in self.cells]
        for syn in self.syns:
            syn.e = 0 * mV
        # connect cell i to cell (i + 1) % num
        pc = h.ParallelContext()
        self.ncs = []
        for i in range(num):
            nc = pc.gid_connect(i, self.syns[(i + 1) % num])
            nc.delay = 1 * ms
            nc.weight[0] = 0.5
            self.ncs.append(nc)
```

```
n = Network(20)
```

Storing spike times

- With 20 cells, it is hard to distinguish the cells when simultaneously plotting the membrane potentials. Let's just store the spike times.
- We begin by modifying
Pyramidal._register_netcon:

```
def _register_netcon(self):  
    self.nc = h.NetCon(  
        self.soma[0](0.5)._ref_v,  
        None,  
        sec=self.soma[0])  
    pc = h.ParallelContext()  
    pc.set_gid2node(self._gid, pc.id())  
    pc.cell(self._gid, self.nc)  
    self.spike_times = h.Vector()  
    self.nc.record(self.spike_times)
```

- When the simulation is over, we can print out the spike times:

```
for i, cell in enumerate(n.cells):  
    print(f'{i}: {list(cell.spike_times)}')
```

- Beginning of output:

```
0: [5.525000000100045, 54.825000000103906]  
1: [8.025000000100079, 57.30000000010447]  
2: [10.525000000099936, 59.77500000010503]  
3: [13.000000000099796, 62.250000000105594]
```

Storing spike times in JSON

- To store spike times in JSON, we can use the following code:

```
import json

with open('output.json', 'w') as f:
    f.write(
        json.dumps({
            i: list(cell.spike_times)
            for i, cell in enumerate(n.cells)},
        indent=4))
```

- JSON is a standard format for data interchange. Libraries are available for most programming languages.

- This creates a file `output.json` which begins:

```
{
    "0": [
        5.525000000100045,
        54.825000000103906
    ],
    "1": [
        8.025000000100079,
        57.30000000010447
    ],
}
```

Raster plots

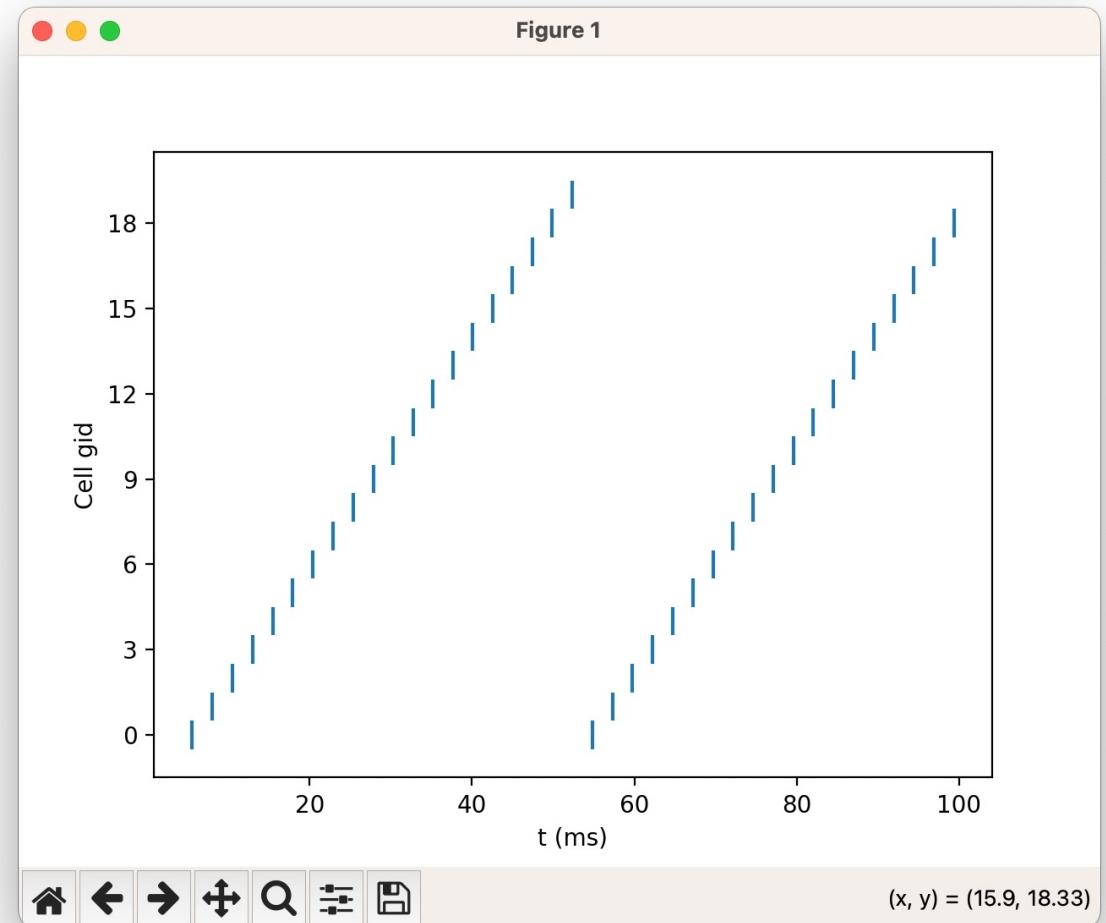
```
from matplotlib.ticker import MaxNLocator

for i, cell in enumerate(n.cells):
    pyplot.vlines(
        list(cell.spike_times), i - 0.5, i + 0.5)

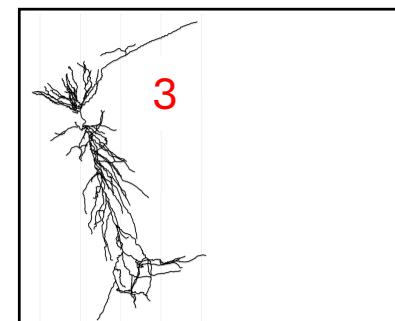
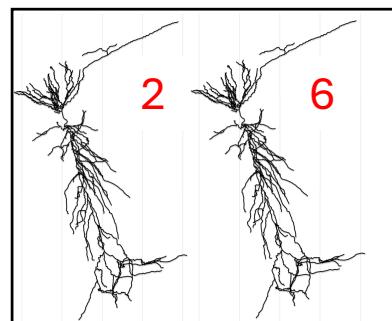
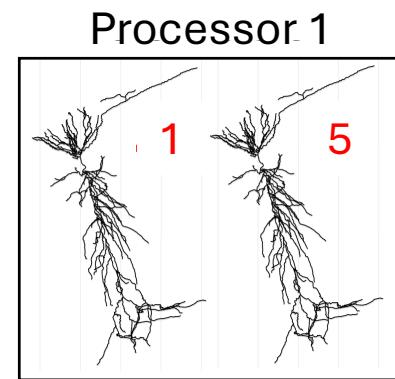
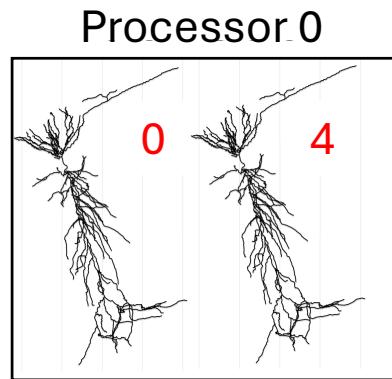
# gids must be integers
pyplot.gca().yaxis.set_major_locator(
    MaxNLocator(integer=True))

pyplot.xlabel('t (ms)')
pyplot.ylabel('Cell gid')

pyplot.show()
```



Simple load-balancing: round-robin



An efficient way to distribute, especially if all cells similar:

```
for gid in range(pc.id(), ncell, pc.nhost()):  
    pc.set_gid2node(gid, pc.id())  
    ...
```

Note: The body is executed at most $\lceil ncell / nhost \rceil$ times in a given process; not $ncell$ times.

Advanced load-balancing: balance the *work* not the cells

- Distribute cells round-robin to all processors, instantiate them.
- Compute an estimate of the computational complexity:

```
def complexity(self):  
    h.load_file('loadbal.hoc')  
    lb = h.LoadBalance()  
    return lb.cell_complexity(sec=self.all[0])
```

- Destroy the cells, send the gid-complexity data to node 0.
- (On node 0): distribute gids such that the next gid goes to the node with the least amount of complexity.
- Send the gids to the nodes; instantiate the cells.

Parallelizing our ring network with round-robin

Very few changes are necessary.

MPI must be initialized before we can use it:

```
h.nrnmpi_init()
```

The Network class only instantiates gids on the current processor.

```
class Network:
    def __init__(self, num):
        pc = h.ParallelContext()
        mygids = list(range(pc.id(), num, pc.nhost()))
        self.cells = [Pyramidal(i) for i in mygids]
        # setup an exciteable ExpSyn on each cell's dendrites
        self.syns = [h.ExpSyn(cell.dend[0](0.5)) for cell in self.cells]
        for syn in self.syns:
            syn.e = 0 * mV
        # connect cell (i - 1) % num to cell i
        self.ncs = []
        for i, syn in zip(mygids, self.syns):
            nc = pc.gid_connect((i - 1) % num, syn)
            nc.delay = 1 * ms
            nc.weight[0] = 1
            self.ncs.append(nc)
```

Parallelizing our ring network

We must modify the initial netstim to ensure it only attaches to gid 0 not to the 0th cell in each process.

```
# drive the 0th cell
if pc.gid_exists(0):
    stim = h.NetStim()
    stim.number = 1
    stim.start = 3
    ncstim = h.NetCon(stim, n.syns[0])
    ncstim.delay = 1
    ncstim.weight[0] = 1
```

Finally, we modify the write to do it on a per-processor basis:

```
with open(f'output{pc.id()}.json', 'w') as f:
    f.write(json.dumps({cell._gid: list(cell.spike_times) for cell in n.cells}, indent=4))
```

Optional: use `pc.py_alltoall` to send all spikes to node 0

```
local_data = {cell._gid: list(cell.spike_times) for cell in n.cells}
all_data = pc.py_alltoall([local_data] + [None] * (pc.nhost() - 1))

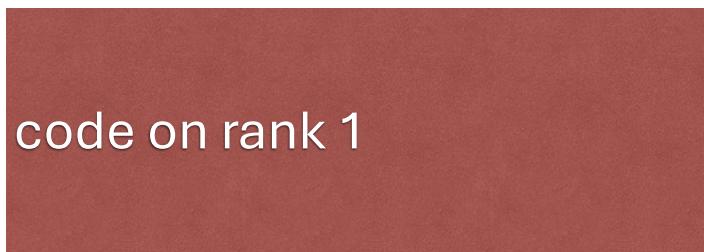
if pc.id() == 0:
    # only do output from node 0
    import json
    combined_data = {}
    for node_data in all_data:
        combined_data.update(node_data)
    with open('output.json', 'w') as f:
        f.write(json.dumps(combined_data, indent=4))
```

Behind the scenes

Automatically exchange any pickleable data type, but there is overhead.



`pc.py_alltoall`, etc

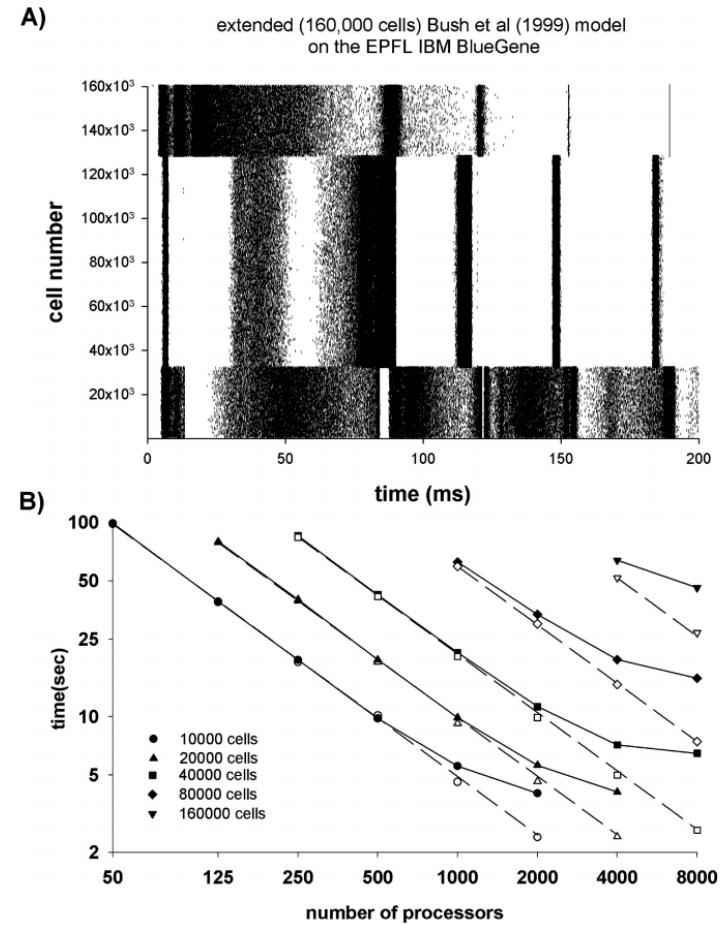
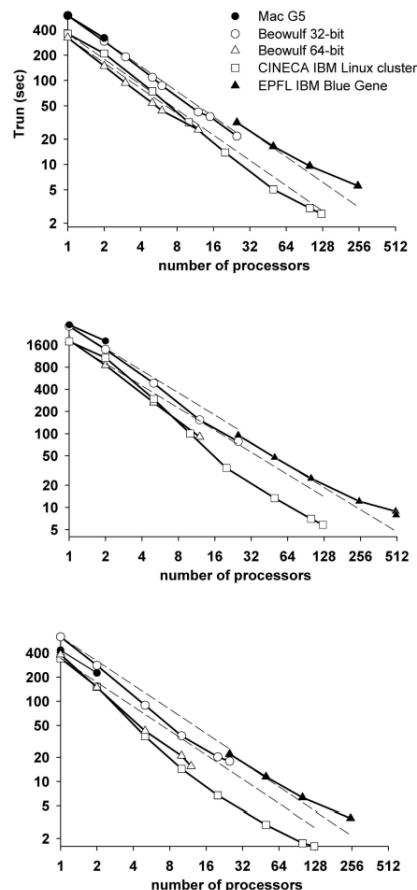
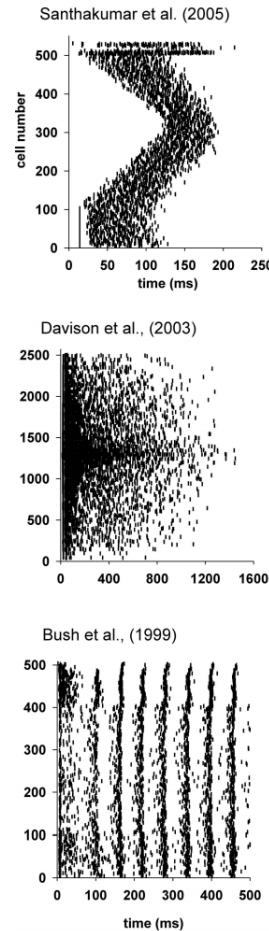


Behind the scenes

Automatically exchange any pickleable data type, but there is overhead.



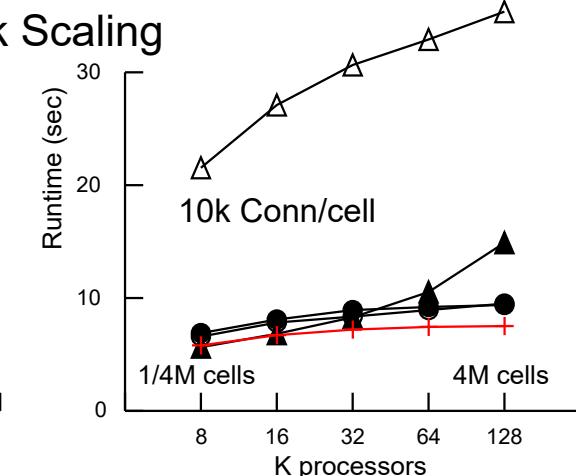
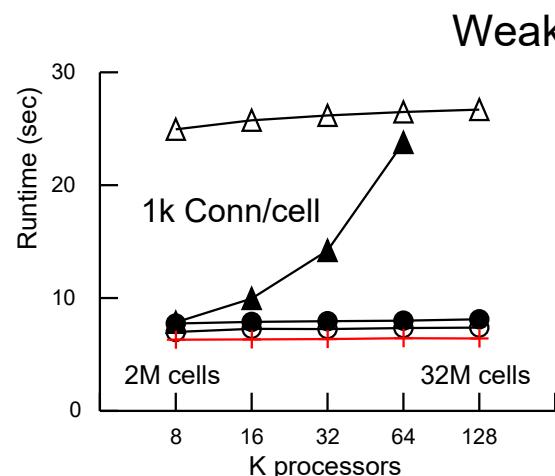
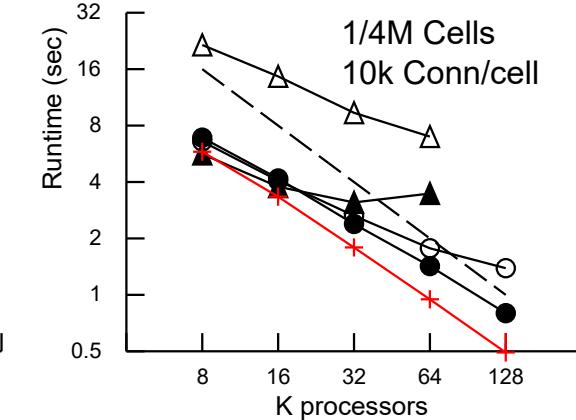
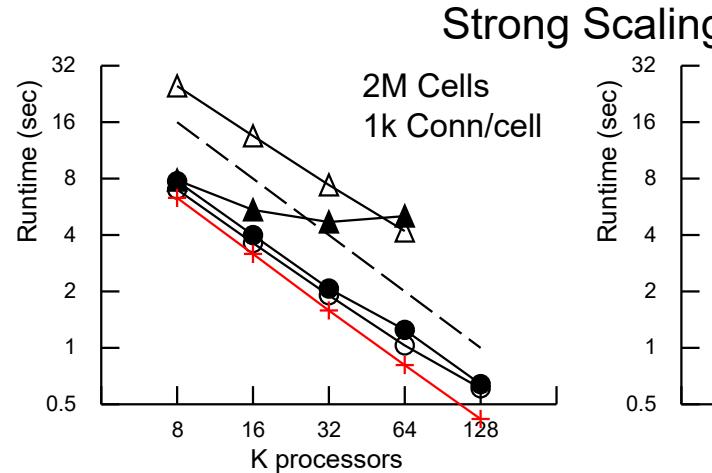
Performance: MPI scaling



Performance: Spike exchange strategies

- △ MPI_ISend – Two Phase, Two Subinterval
- ▲ Allgather
- DCMF_Multicast – Two Phase, Two Subinterval
- Record-Replay – One Subinterval
- + Computation Time (includes queue)

Artificial Spiking Net
Blue Gene/P
Argonne National Lab



Tip

- In large network models, we expect an average of more than one spike per 0.025 ms (NEURON's default timestep).
- If we use a variable step solver, we will capture those times correctly, but this will require restarting the integrator frequently.
- In practice, it is usually better to use a fixed step solver for network models.



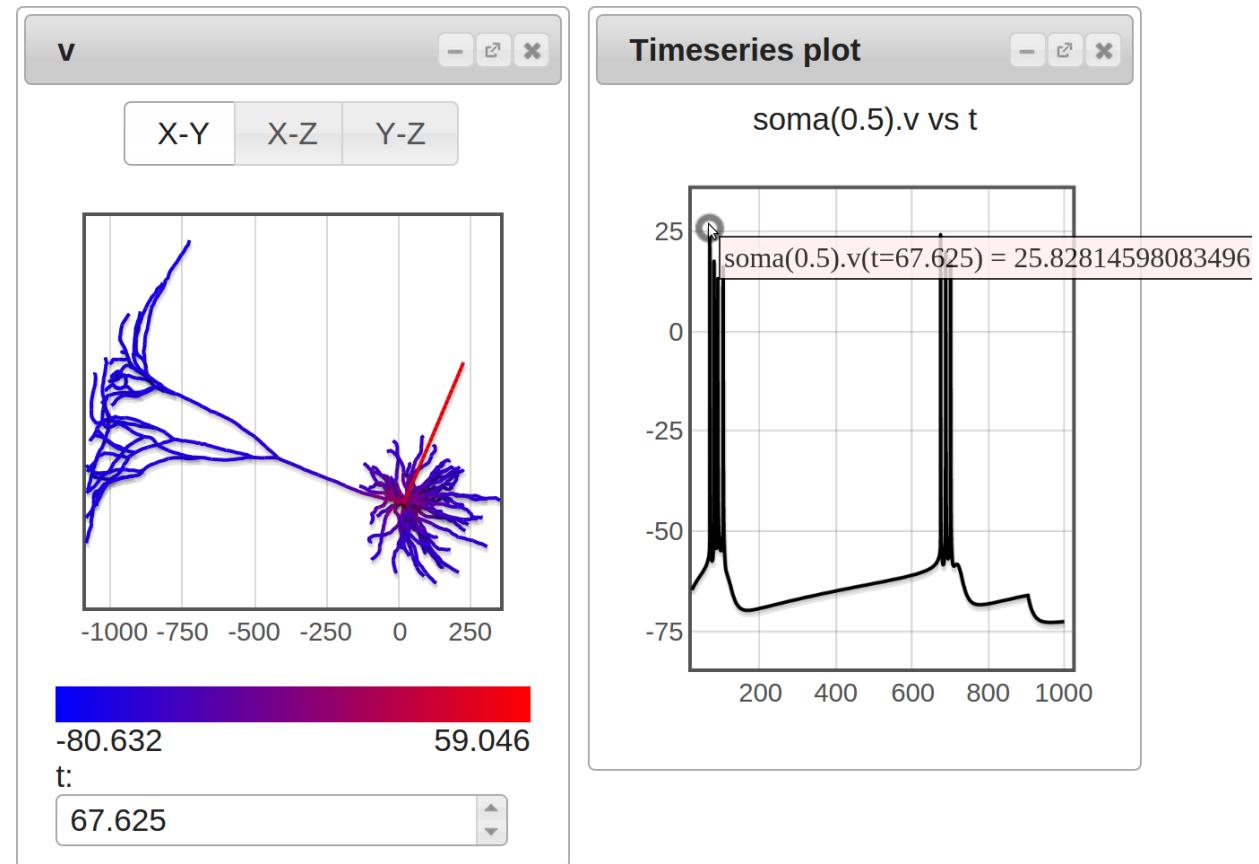
Question

- Suppose we ran a large network simulation and stored the spike times.
- Suppose further we now realize we want to know the time series of the m variable in the center of the soma of cell 5.
- Do we have to modify our code to store that variable and rerun the entire simulation?

Tip: Store synaptic events and recreate cells as needed



initial conditions
+
synaptic events \rightarrow neuron dynamics



Using spike data to recreate a variable of interest

We will need `vecevent.mod`. If you have NEURON, this file should be on your computer somewhere. Alternatively, you can download it from:

<https://github.com/neuronsimulator/nrn/blob/master/share/examples/nrniv/netcon/vecevent.mod>

Using spike data to recreate a variable of interest

```
import json
from neuron import h
from neuron.units import ms, mV
from PyNeuronToolbox import morphology
from matplotlib import pyplot
h.load_file('stdrun.hoc')
num_cells = 20

# class Pyramidal as before

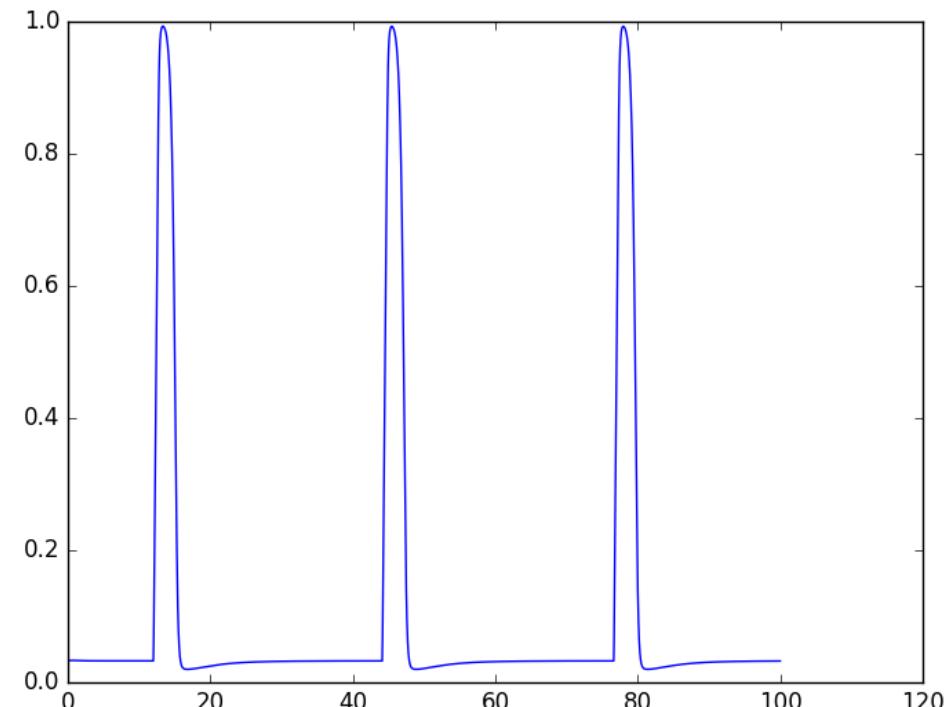
# read spike times
with open('output.json') as f:
    spike_times_by_cell = json.load(f)
```

(continued)

Using spike data to recreate a variable of interest

```
def get_m(gid):
    p = Pyramidal(gid)
    # recreate synaptic inputs (here, only one; you may have multiple)
    precell = (gid - 1) % num_cells
    vs = h.VecStim()
    spike_vec = h.Vector(spike_times_by_cell[str(precell)])
    vs.play(spike_vec)
    syn = h.ExpSyn(p.dend[0](0.5))
    nc = h.NetCon(vs, syn)
    nc.delay = 1 * ms
    nc.weight[0] = 1
    # setup recording
    t = h.Vector().record(h._ref_t)
    m = h.Vector().record(p.soma[0](0.5)._ref_m_hh)
    # do run
    pc = h.ParallelContext()
    pc.set_maxstep(10 * ms)
    h.finitialize(-69 * mV)
    pc.psolve(100 * ms)
    return t, m

t, m = get_m(5)
pyplot.plot(t, m)
pyplot.show()
```



Pseudo-randomness

Ensuring reproducibility while sampling from distributions.

When would we want randomness?

- Stochastic input (e.g., spikes arriving via a Poisson process).
- Stochastic channel gating.
- Random channel parameters.
- Random morphology.
- Random network connections.

Guiding principles

- Randomness should be **reproducible**. e.g., use random seeds.
- Randomness should be **independent of the number of processors**.

Philox/Random123

- Random123 is a **cryptographic quality** generator, suitable for managing separate **independent**, **reproducible**, and **restartable** streams; the algorithm used is called Philox in numpy and in the paper:

Parallel Random Numbers: As Easy as 1, 2, 3

John K. Salmon,^{*} Mark A. Moraes, Ron O. Dror, and David E. Shaw^{*†}
D. E. Shaw Research, New York, NY 10036, USA

```
from neuron import h
r = h.Random()
r.Random123(seed, id1, id2)
print(r.uniform(0,1)) # or binomial, or geometric, or lognormal, or negexp, or ...
```

- A reasonable choice for id1 is the gid of the cell; a reasonable choice for id2 is the section index in cell.all or the synapse index.

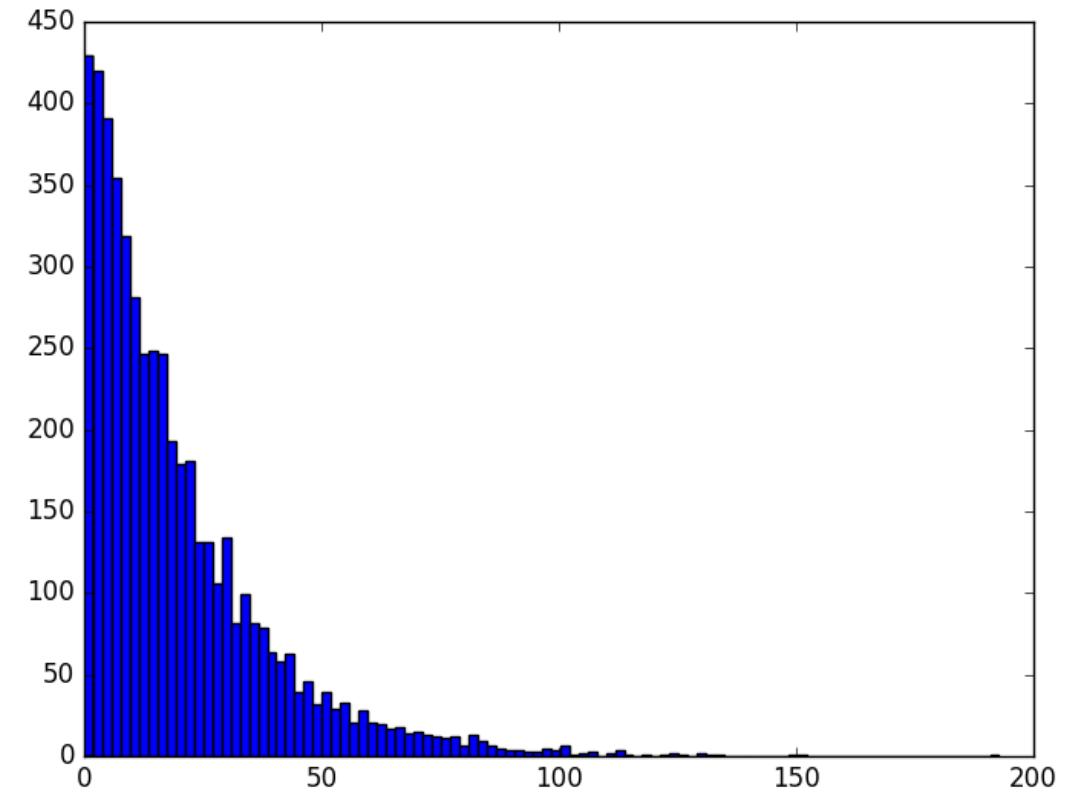
Example: random synaptic inputs (“noise”)

```
from neuron import h, gui
import numpy
from matplotlib import pyplot

r = h.Random()
r.negexp(1)
r.Random123(1, 2, 3)
ns = h.NetStim()
ns.noiseFromRandom(r)
ns.number = 1e20
ns.start = 5
ns.interval = 20
ns.noise = 1

spike_times = h.Vector()
nc = h.NetCon(ns, None)
nc.record(spike_times)
pc = h.ParallelContext()
pc.set_maxstep(10)
h.finitialize(-65)
pc.psolve(100000)

spike_times = spike_times.as_numpy()
isis = spike_times[1:] - spike_times[:-1]
print(numpy.mean(isis), numpy.std(isis))
pyplot.hist(isis, 100)
pyplot.show()
```



Example: random distribution of leak conductance

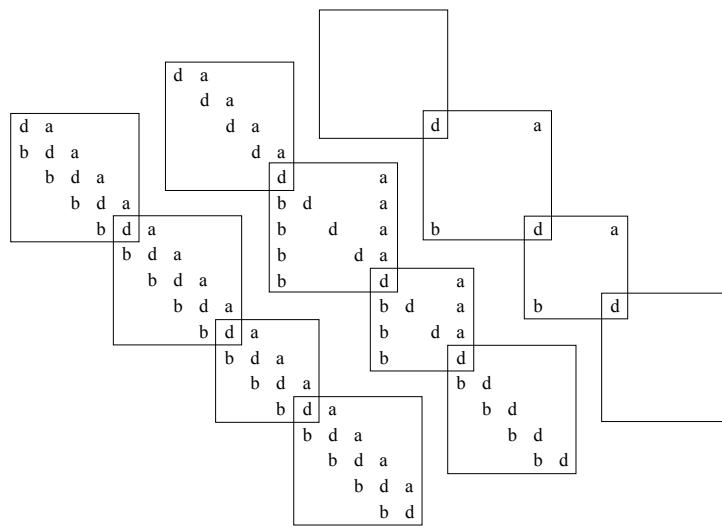
```
import sys
seed = float(sys.argv[1])

class Pyramidal:
    def __init__(self, gid):
        self._gid = gid
        self._setup_morphology()
        self._discretize()
        self.random_stream = [self.new_random_stream(i) for i in range(len(self.all))]
        self._add_channels()
        self._register_netcon()
    def new_random_stream(self, id2):
        r = h.Random()
        r.Random123(seed, self._gid, id2)
        return r
    def _add_channels(self):
        for sec in self.soma:
            sec.insert(h.hh)
        for sec, stream in zip(self.all, self.random_stream):
            sec.insert(h.pas)
            for seg in sec:
                seg.pas.g = max(0, stream.normal(1e-3, 8e-6))
    # the rest stays as before
```

If you try this with seed 1, you'll notice the raster looks significantly different. Why?

Multisplit

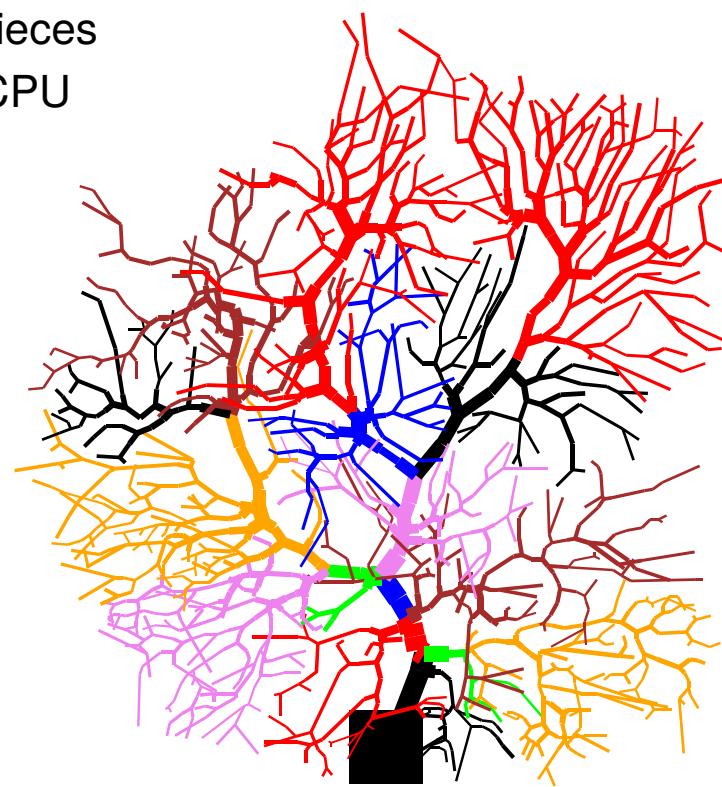
Improve load balancing with multisplit



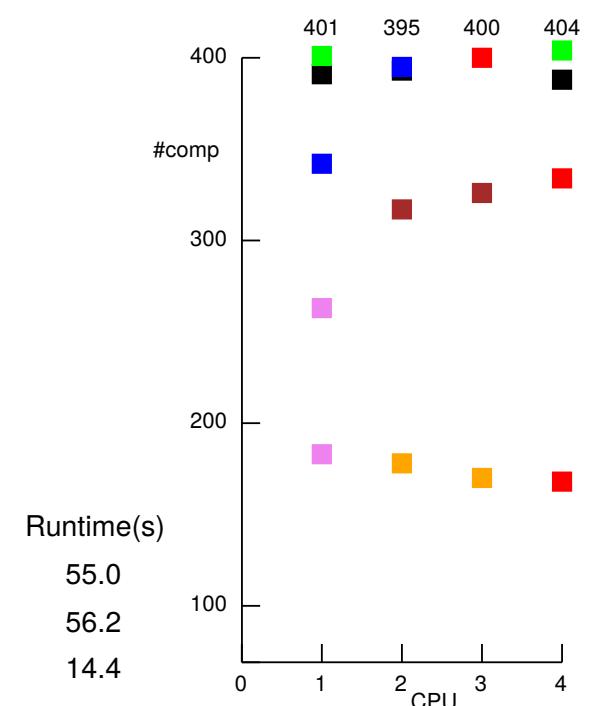
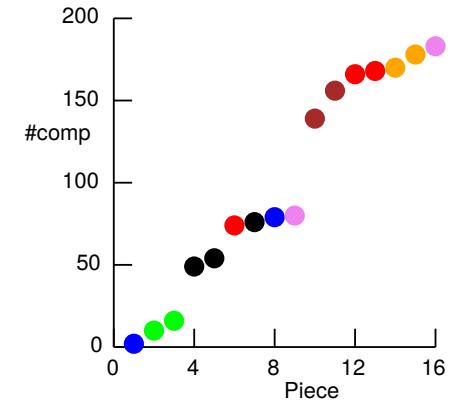
Multisplit algorithm described in Hines et al 2008.

DOI: 10.1007/s10827-008-0087-5

16 Pieces
4 CPU

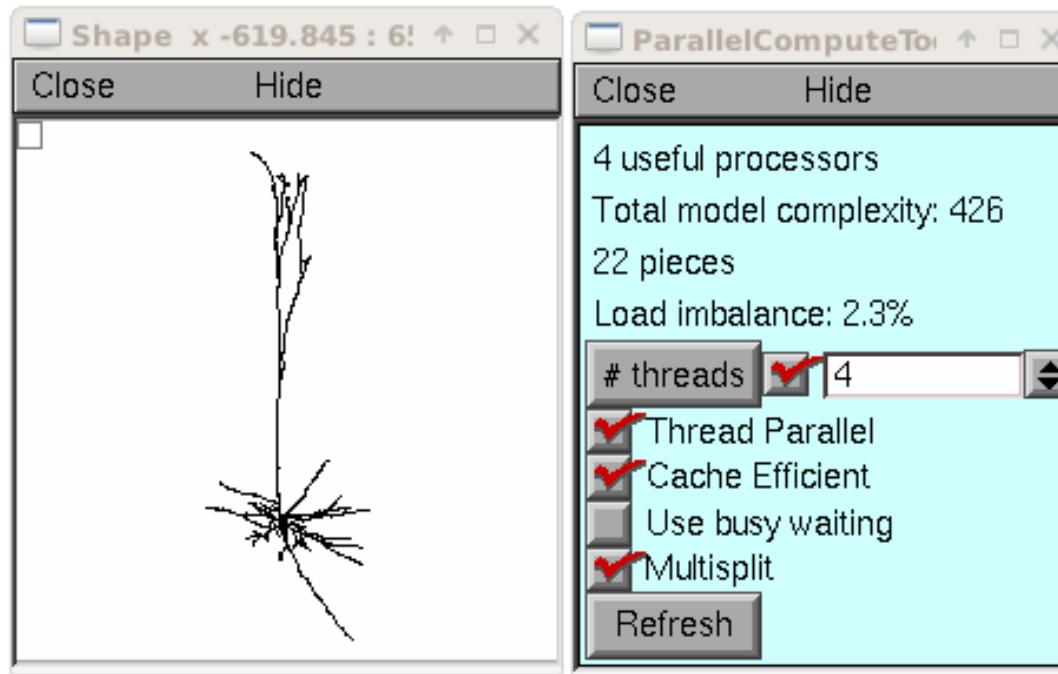


CPU	Computation	Exchange	
0	13.82	0.56	
1	13.35	1.03	16 pieces, 1 cpu
2	13.47	0.90	wholecell, 1 cpu
3	13.56	0.82	16 pieces, 4 cpu



Using multisplit (threads)

When not using MPI, enabling thread-based multisplit is as easy as clicking a checkbox:



Using multisplit (MPI)

For process-based multisplit (with MPI), use `pc.multisplit` to declare split nodes:

```
pc.multisplit(seg, subtreeid)
```

After all split nodes are declared, **every** process must execute:

```
pc.multisplit()
```

If created, destroy any parts of the cell that do not belong on the processor.

Rules:

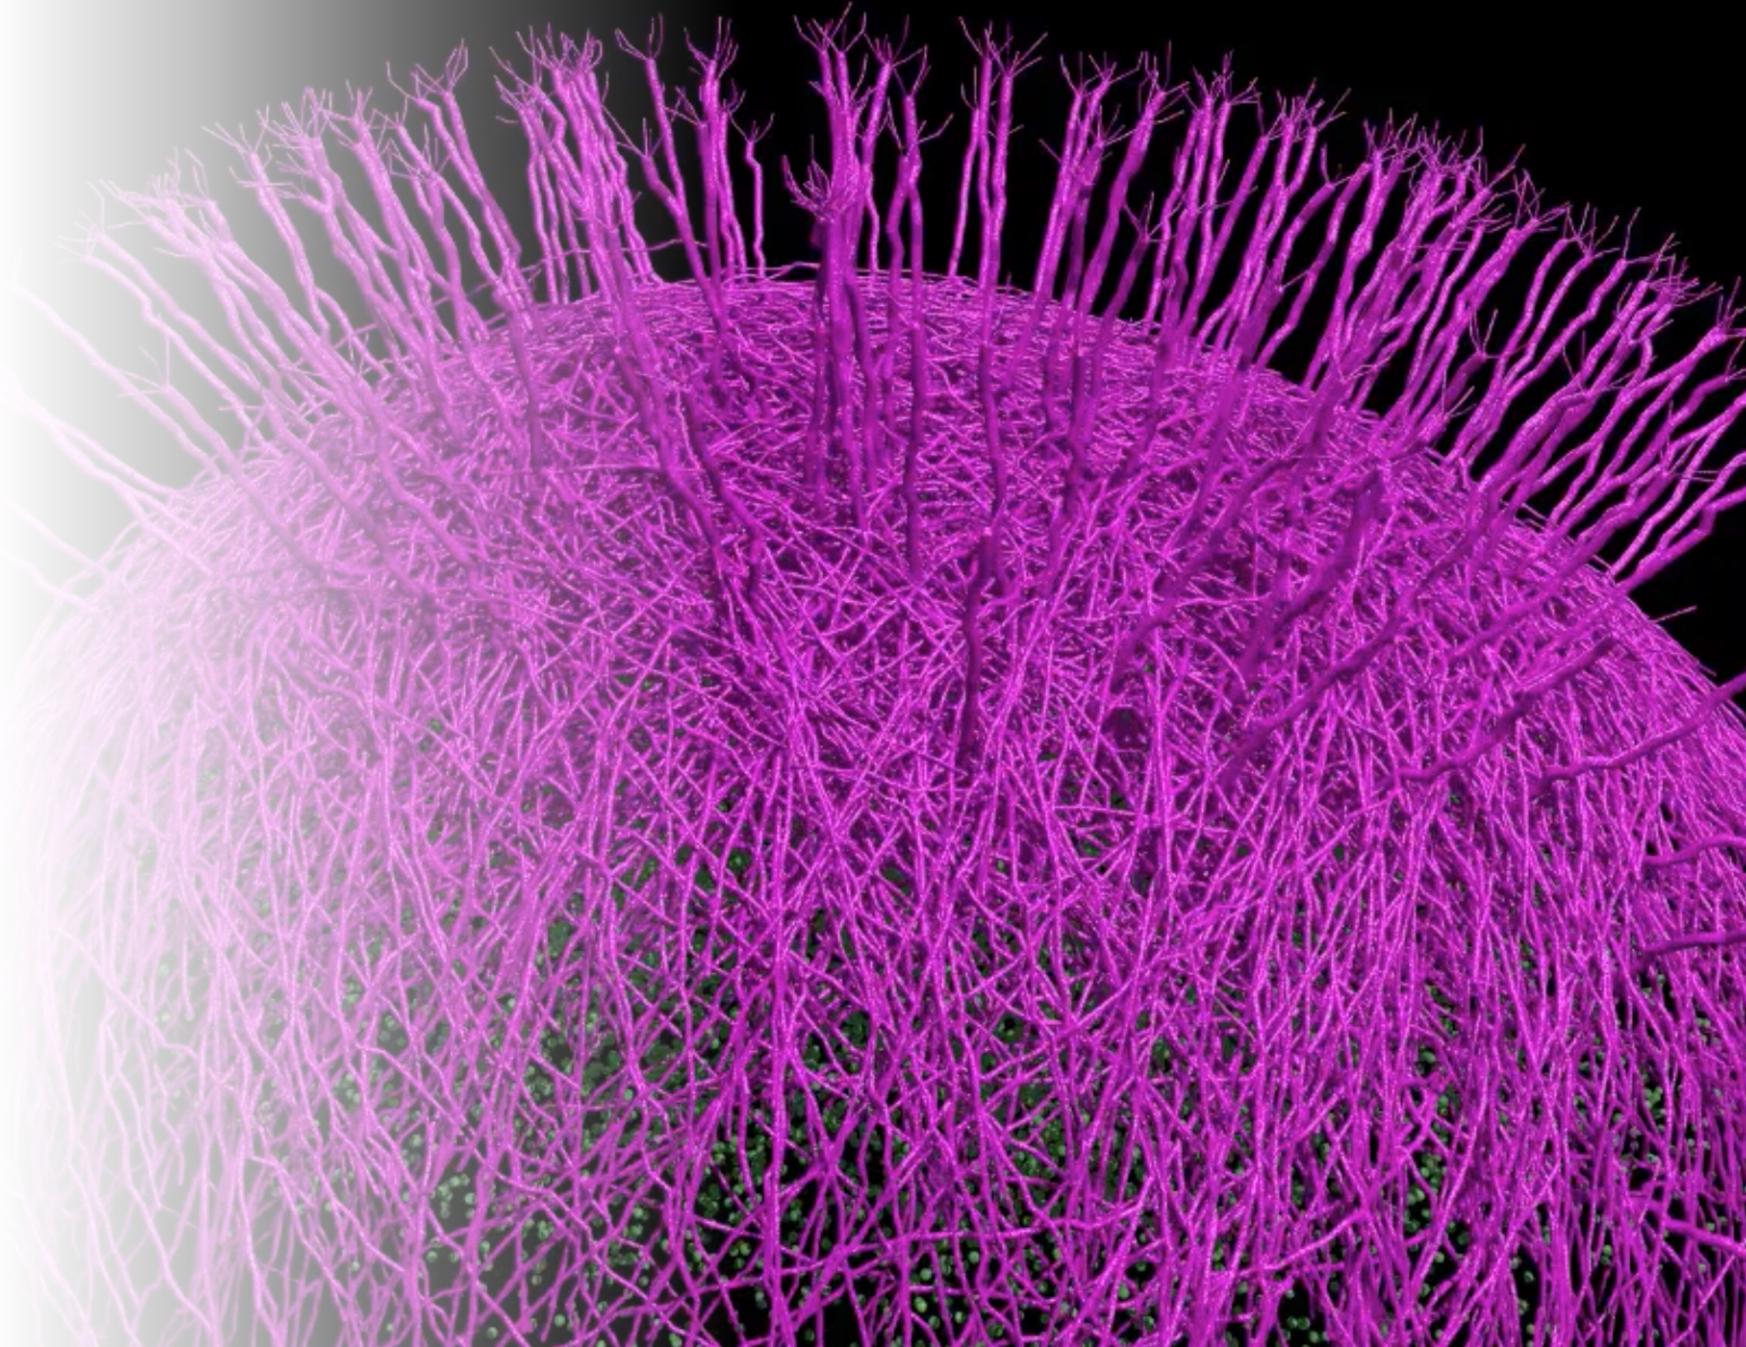
- Each subtree can have at most two split nodes.
- Does not support variable step, linear mechanisms, extracellular, or reaction-diffusion.
- `h.distance` cannot compute path distances that cross a split node.

Tip: For load balancing, it is sometimes convenient to split cells into more pieces than processes.

Example:

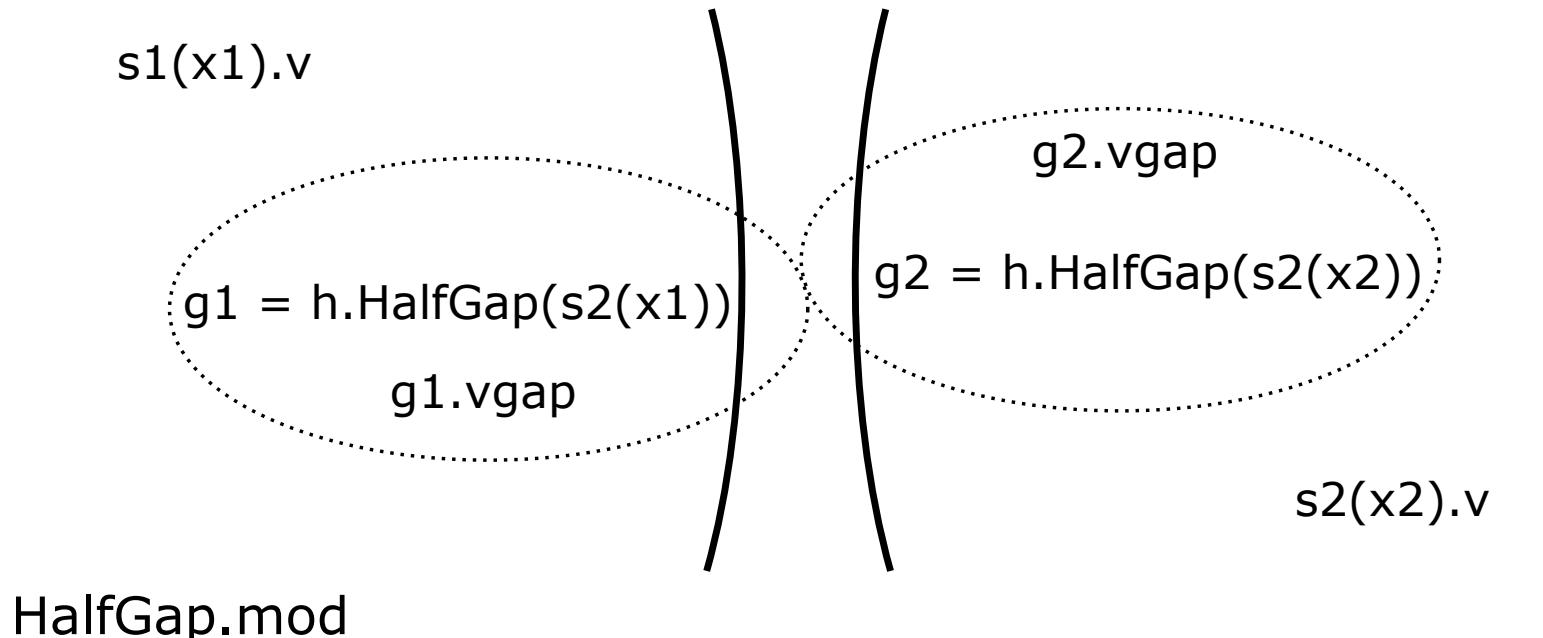
Migliore et al., 2014

- Migliore et al 2014 used multisplit to improve load balancing on a model of the olfactory bulb.
modeldb.science/151681
- See, in particular, the file multisplit_distrib.py.



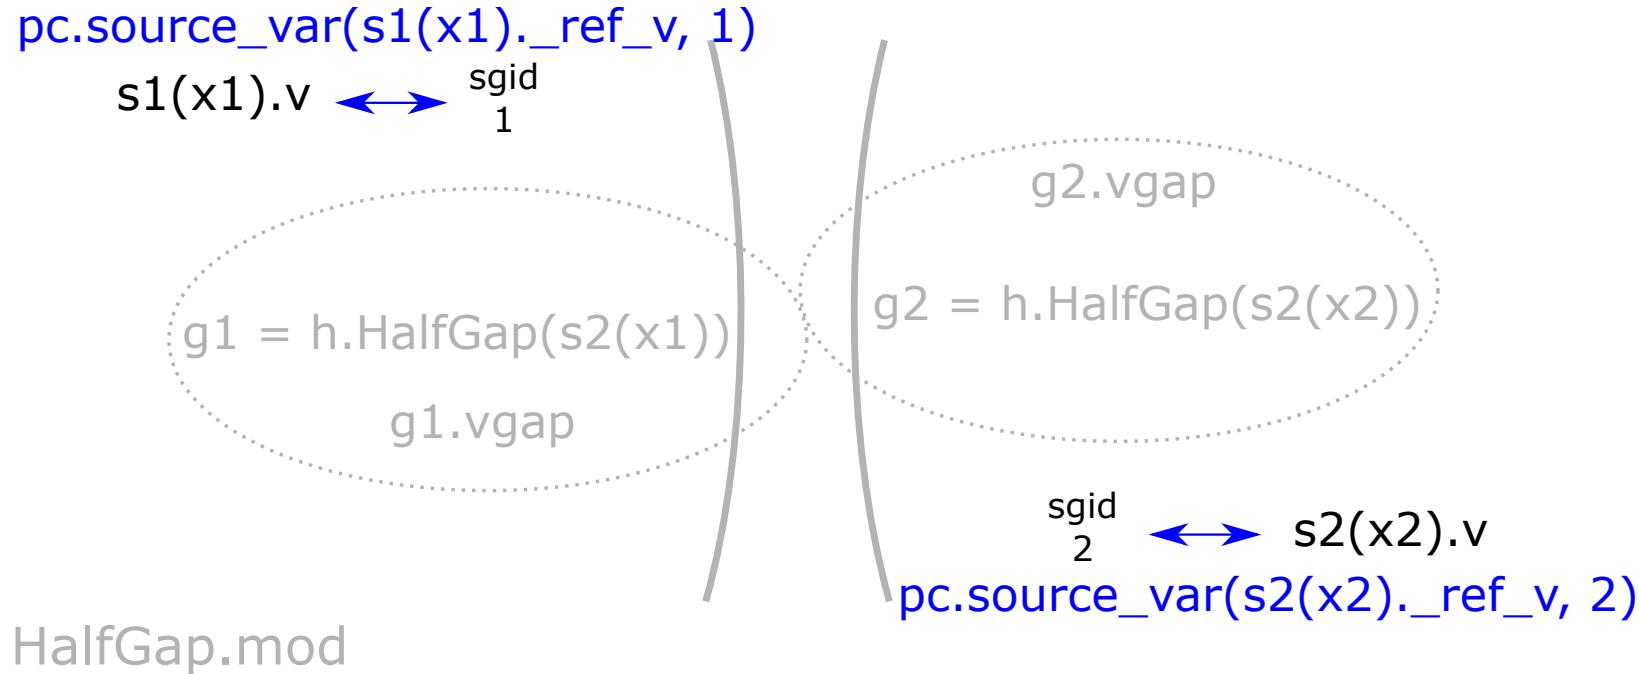
Gap junctions

Continuous voltage exchange



```
NEURON  {
    POINT_PROCESS HalfGap
    ELECTRODE_CURRENT i
    RANGE r, i, vgap
}
PARAMETER { r = 1e9 (megohm) }
ASSIGNED {
    v (millivolt)
    vgap (millivolt)
    i (nanoamp)
}
CURRENT { i = (vgap - v) / r }
```

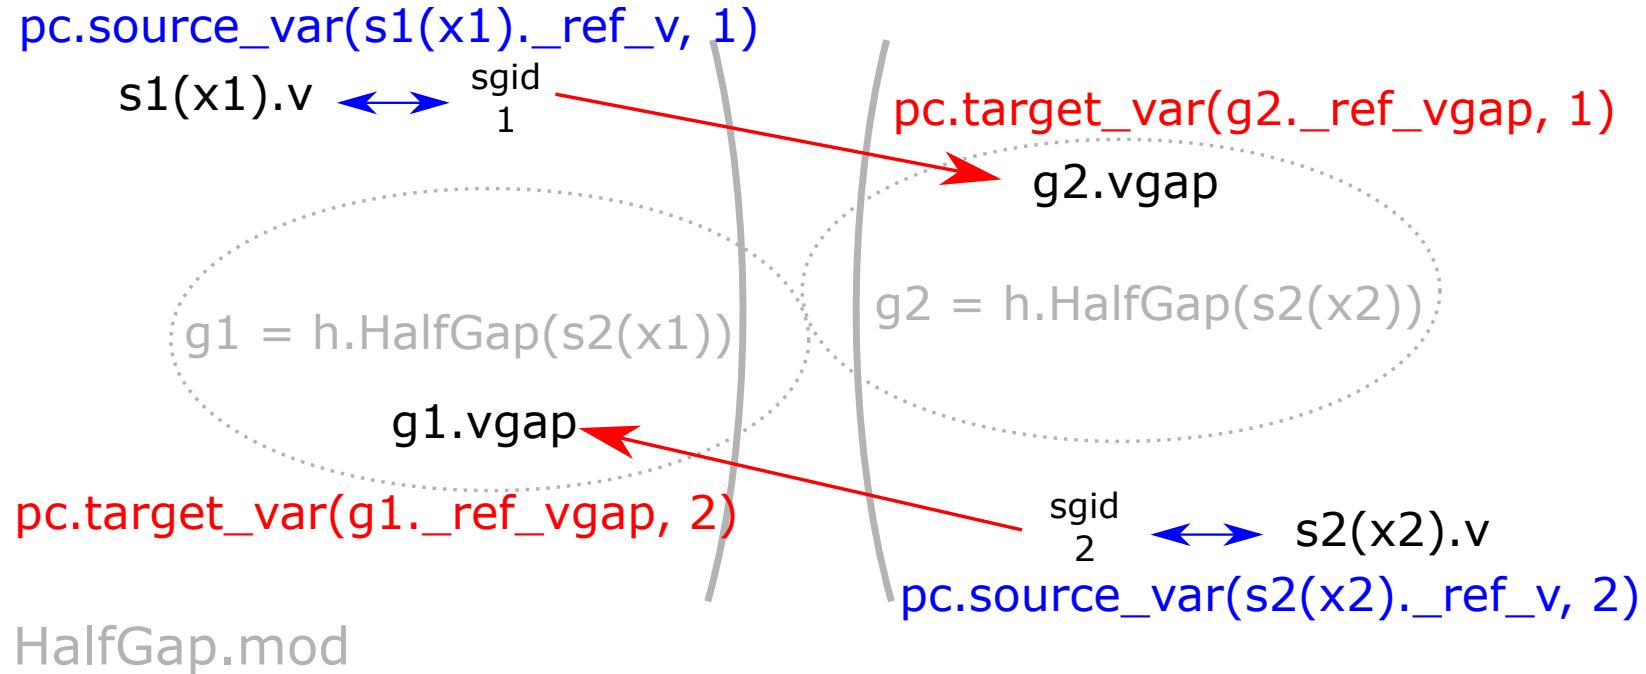
pc.source_var to declare source sgid



```
NEURON {  
    POINT_PROCESS HalfGap  
    ELECTRODE_CURRENT i  
    RANGE r, i, vgap  
}  
PARAMETER { r = 1e9 (megohm) }
```

```
ASSIGNED {  
    v (millivolt)  
    vgap (millivolt)  
    i (nanoamp)  
}  
CURRENT { i = (vgap - v) / r }
```

pc.target_var to declare target connection



```
NEURON {  
    POINT_PROCESS HalfGap  
    ELECTRODE_CURRENT i  
    RANGE r, i, vgap  
}  
PARAMETER { r = 1e9 (megohm) }
```

```
ASSIGNED {  
    v (millivolt)  
    vgap (millivolt)  
    i (nanoamp)  
}  
CURRENT { i = (vgap - v) / r }
```

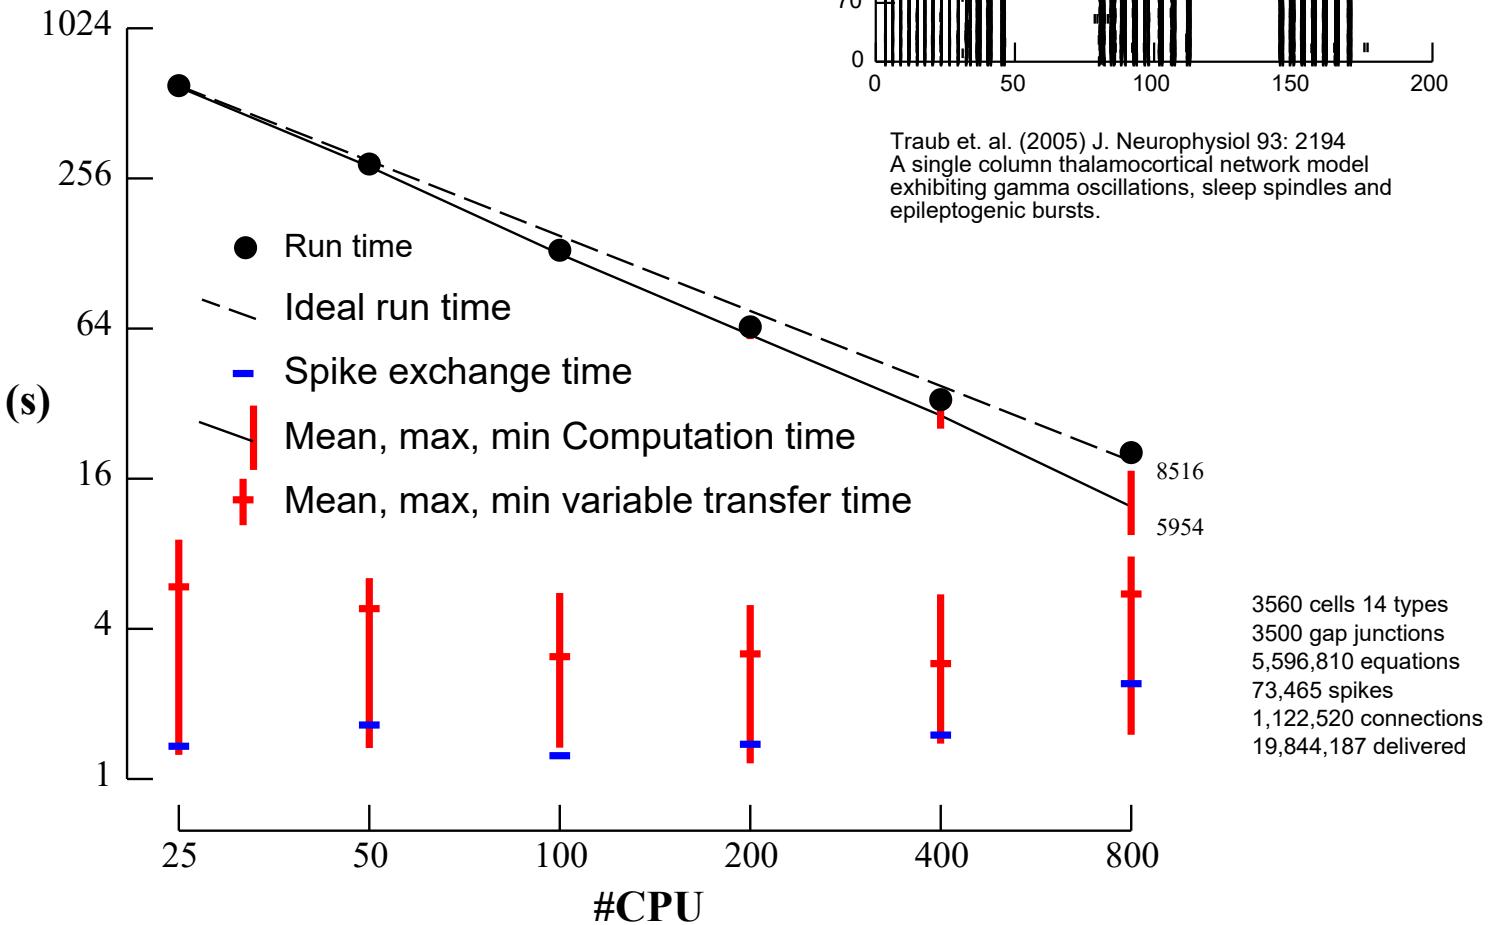
Performance: Traub model

Pittsburgh Supercomputing Center

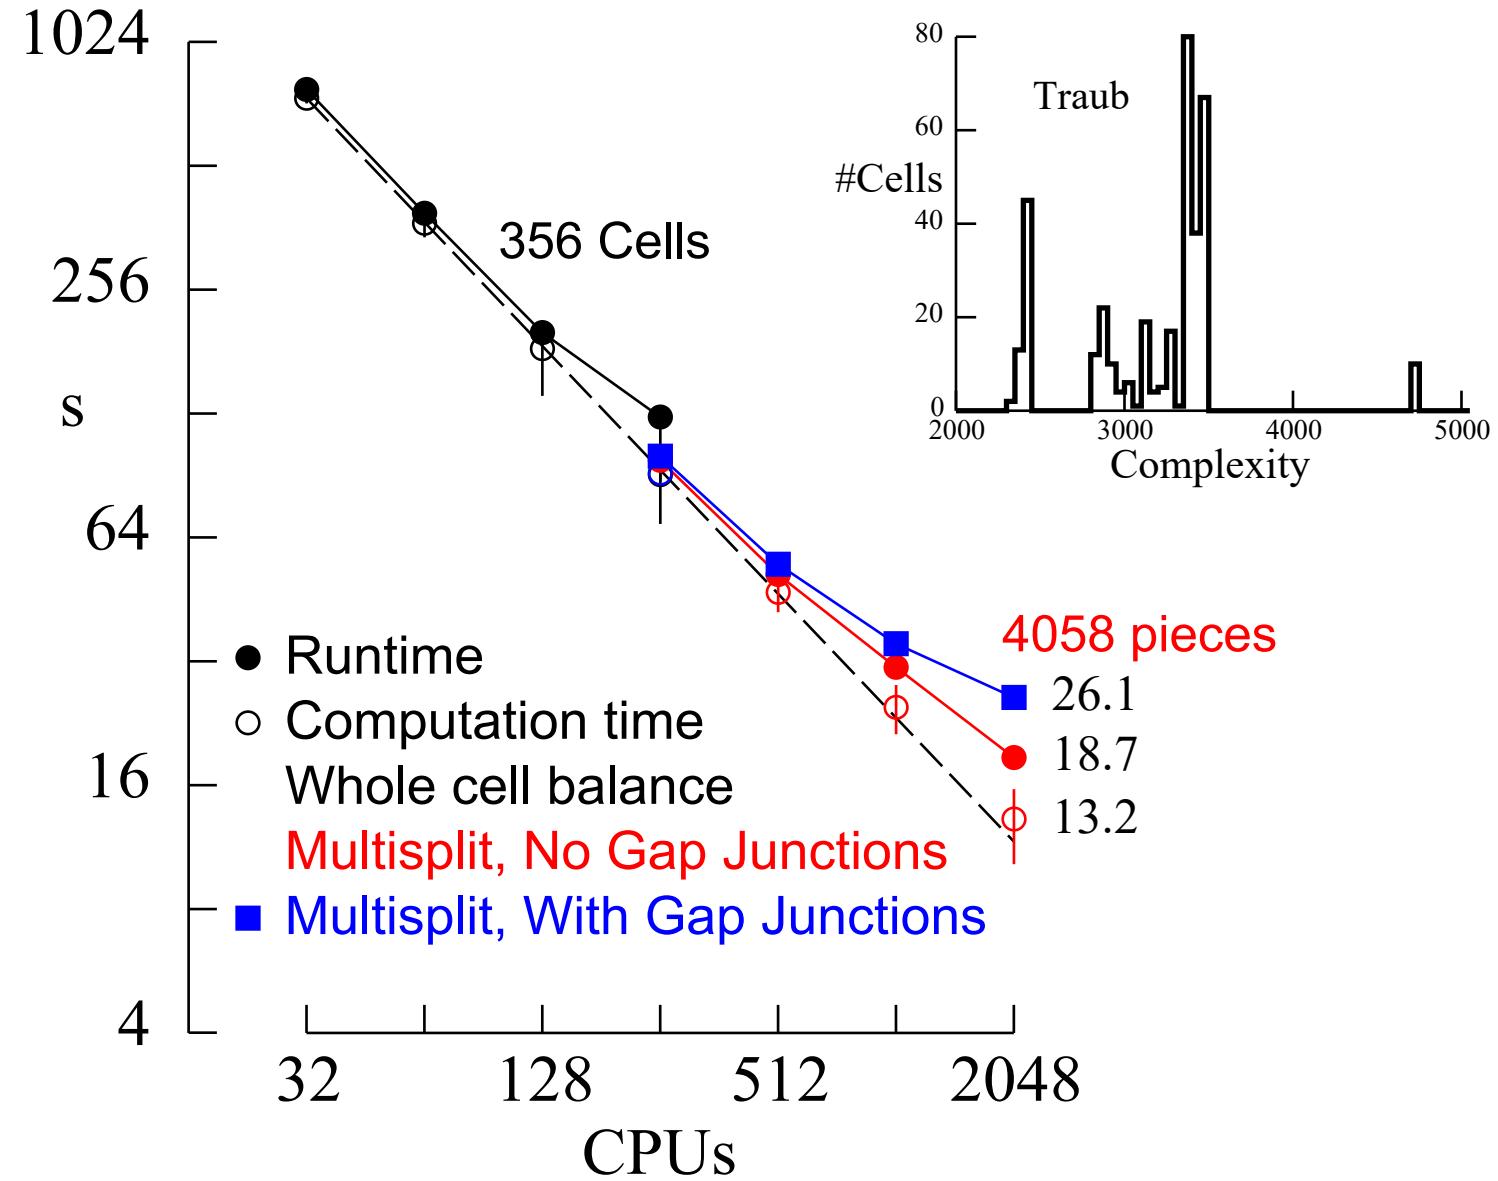
Bigben

Cray XT3

2068 2.4 GHz Opteron Processors



Performance: Traub model with multisplit



Finally: Subworlds

Use `pc.subworlds` to combine parallel simulation with parallel bulletin-board based parameter search.

```
from neuron import h
h.nrnmpi_init()
pc = h.ParallelContext()
pc.subworlds(2)

from model import runmodel

pc.runworker()

for ncell in range(5, 10):
    pc.submit(runmodel, ncell, 1, 100)

while (pc.working()):
    print(pc.pyret())

pc.done()
h.quit()
```

Note: Unless memory on a single node is a limiting factor, you will likely want either 1 subworld (everything) or `pc.nhost()` subworlds. In the first case, there is no need to use subworlds since simulations are run one at a time; in the other extreme, there is also no need since each simulation runs on a single processor.